

# mojaloop

## PI-9 Performance Workstream April 2020

**Team (Alphabetical Order):**

- Coil
- Confluent
- CrossLake
- ModusBox
- Sybrin

mojaloop

# mojaloop

## Performance



MODUSBOX

### Team (Alphabetical Order):

Bryan Schneider  
James Bush  
Miguel de Barros (Presenter)  
Rajiv Mothilal  
Roman Pietrzak  
Sam Kummary  
Valentin Genev

# Performance – PI-8 Performance (Characterization) Report

1. Financial Transaction End-to-End
2. Financial Transaction Prepare-only
3. Financial Transaction Fulfil-only
4. Individual Mojaloop Component Characterization
  - a. Mojaloop Services & Handlers
  - b. Mojaloop Streaming Architecture & Libraries
  - c. Database

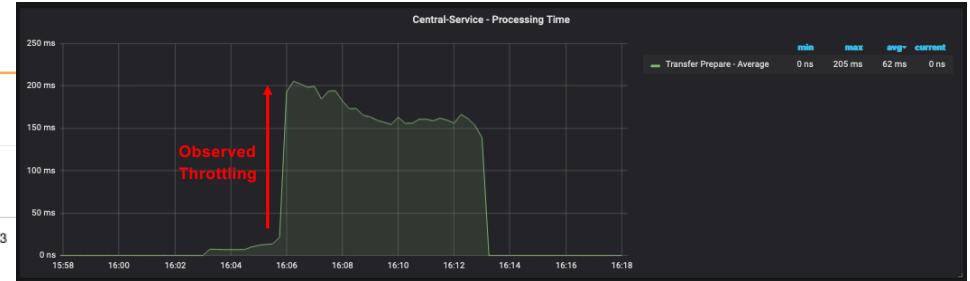
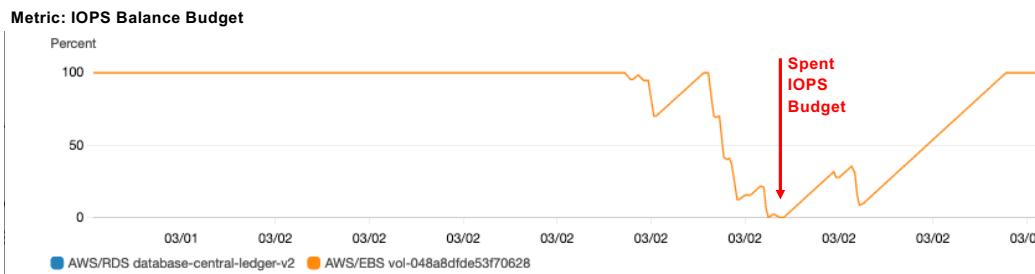
Download Report URI:

[https://github.com/mojaloop/documentation-artifacts/blob/master/presentations/April%202020%20Community%20Event/Pre\\_Read/202003-PI8-Performance-Report-v1.0.pdf](https://github.com/mojaloop/documentation-artifacts/blob/master/presentations/April%202020%20Community%20Event/Pre_Read/202003-PI8-Performance-Report-v1.0.pdf)

# Resolved Limiting Factors on AWS cloud-provider

## 1. IOPS Budget Balance Throttling

- a) Default 300 IOPS limit for GP2 instances @ 100GB
- b) Avoid by Provisioning dedicated IOPS

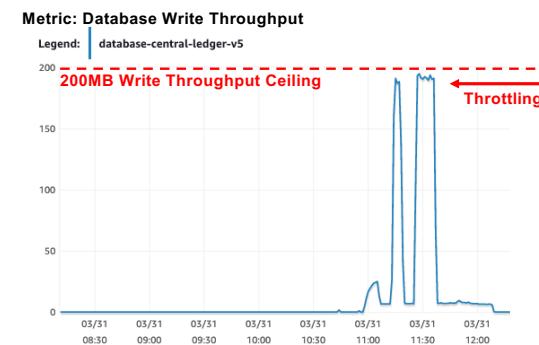


## 2. IOPS Ceiling

- a) Ensure enough IOPS provisioned on storage
- b) Performance tests used 4k—20k IOPS (depends on expected Write Throughput)

## 3. Write Throughput Ceiling

- a) Ensure dedicated network throughput for attached storage matches write throughput requirements
  - a) 100MB/s - db.r4.xlarge
  - b) 200MB/s - db.r4.2xlarge
  - c) 375MB/s - db.r4.2xlarge



Key:  
● Migrated to Master branch (v9.5.x)  
● Performance branch only

# Performance – Enhancement Changes

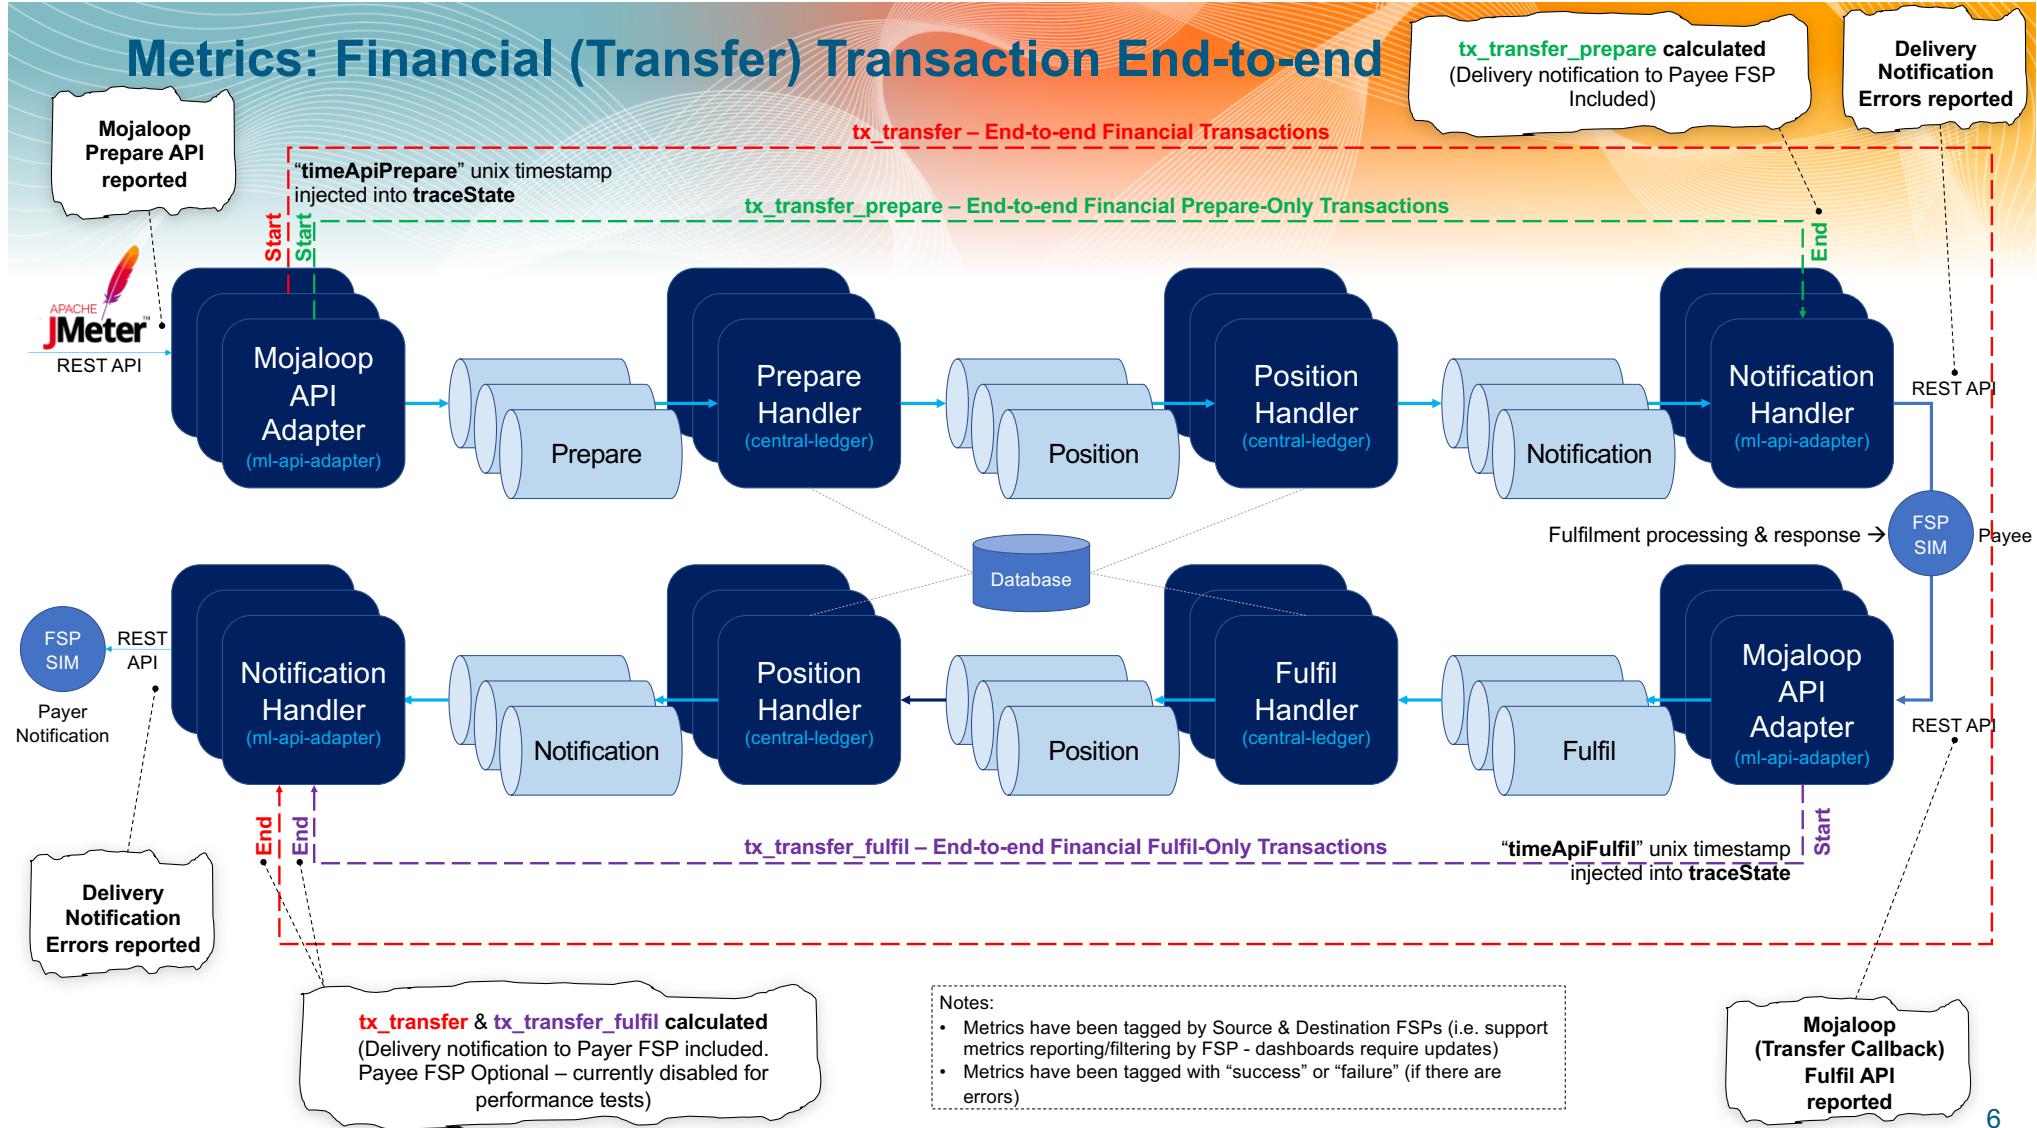
1. Optimized Logging (central-services-logger)
    - a) Added conditional logging to reduce unnecessary processing/parsing logic:
      - ● ML-API-Adapter
      - ● Central-Ledger
      - ● Simulator
      - ● Central-Services-Shared
      - ● Central-Services-Stream
- ```

-   Logger.debug('create::payload(%s)', JSON.stringify(request.payload))
+   Logger.isDebugEnabled && Logger.debug('create::payload(%s)', JSON.stringify(request.payload))

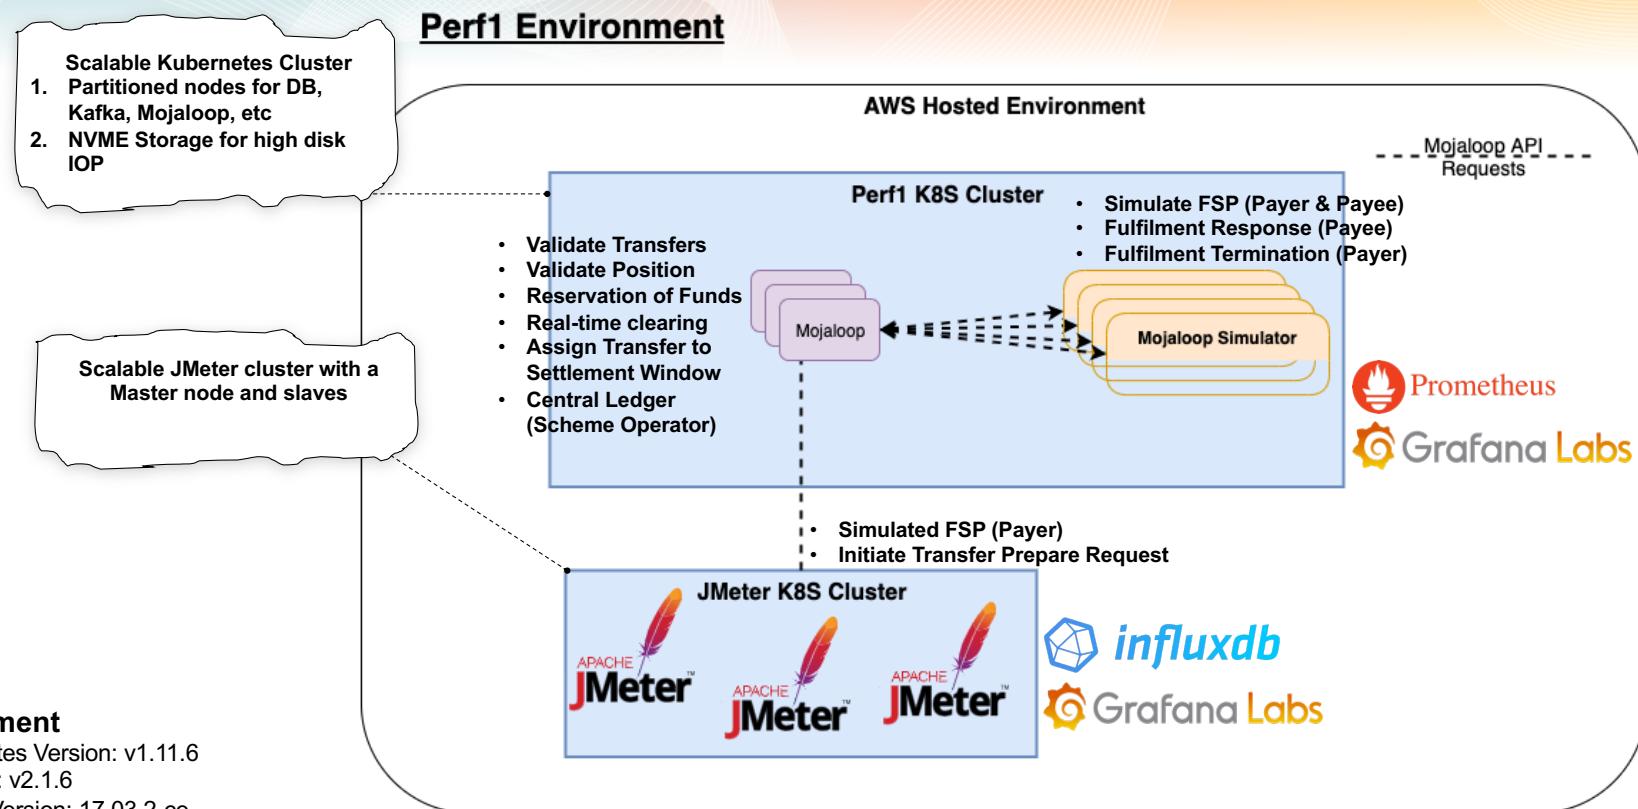
```
2. Caching of Non-Stateful Data
    - a) ● Participants - central-ledger
    - b) ● Participant Accounts - central-ledger
    - c) ● Participant Limits - central-ledger
3. Optimized SQL Queries
    - a) Duplicate Checks
      - a) ● Insert-only logic
    - b) Position changes
      - a) ● Reduced locks by in-line calculation

4. Metrics (central-services-metrics)
  - a) v9.5.x Release supports Histograms, and Summary (new)
  - b) New Instrumentations
    - ML-API-Adapter - reported by Notification Handler
      - ● Notification Delivery
      - ● Errors for Notification Delivery
      - ● End-to-end Financial Transactions - 'tx\_transfer'
      - ● Prepare-only Financial Transactions - 'tx\_transfer\_prepare'
      - ● Fulfil-only Financial Transactions - 'tx\_transfer\_fulfil'
      - ● Transaction processing time buckets
      - ● % of Transaction longer than 1 sec
    - Central-Ledger – reported by Prepare, Position & Fulfil Handlers
      - ● Domain (functional) logic
      - ● SQL queries, inserts & updates
      - ● Cache Hits / Misses
5. Tracing (event-sdk)
  - a) v9.5.x Release supports base64 encoded traceState - containing spanId & custom key-value pairs
    - ● timeApiPrepare – unix timestamp for start of Prepare Process
    - ● timeApiFulfil – unix timestamp for start of Fulfil Process
6. Optimized Central-Ledger Handlers → Proof of Concept
  - a) ● Prepare + Position -> Prepare-Position Handler 5
  - b) ● Fulfil + Position -> Fulfil-Position Handler

## Metrics: Financial (Transfer) Transaction End-to-end



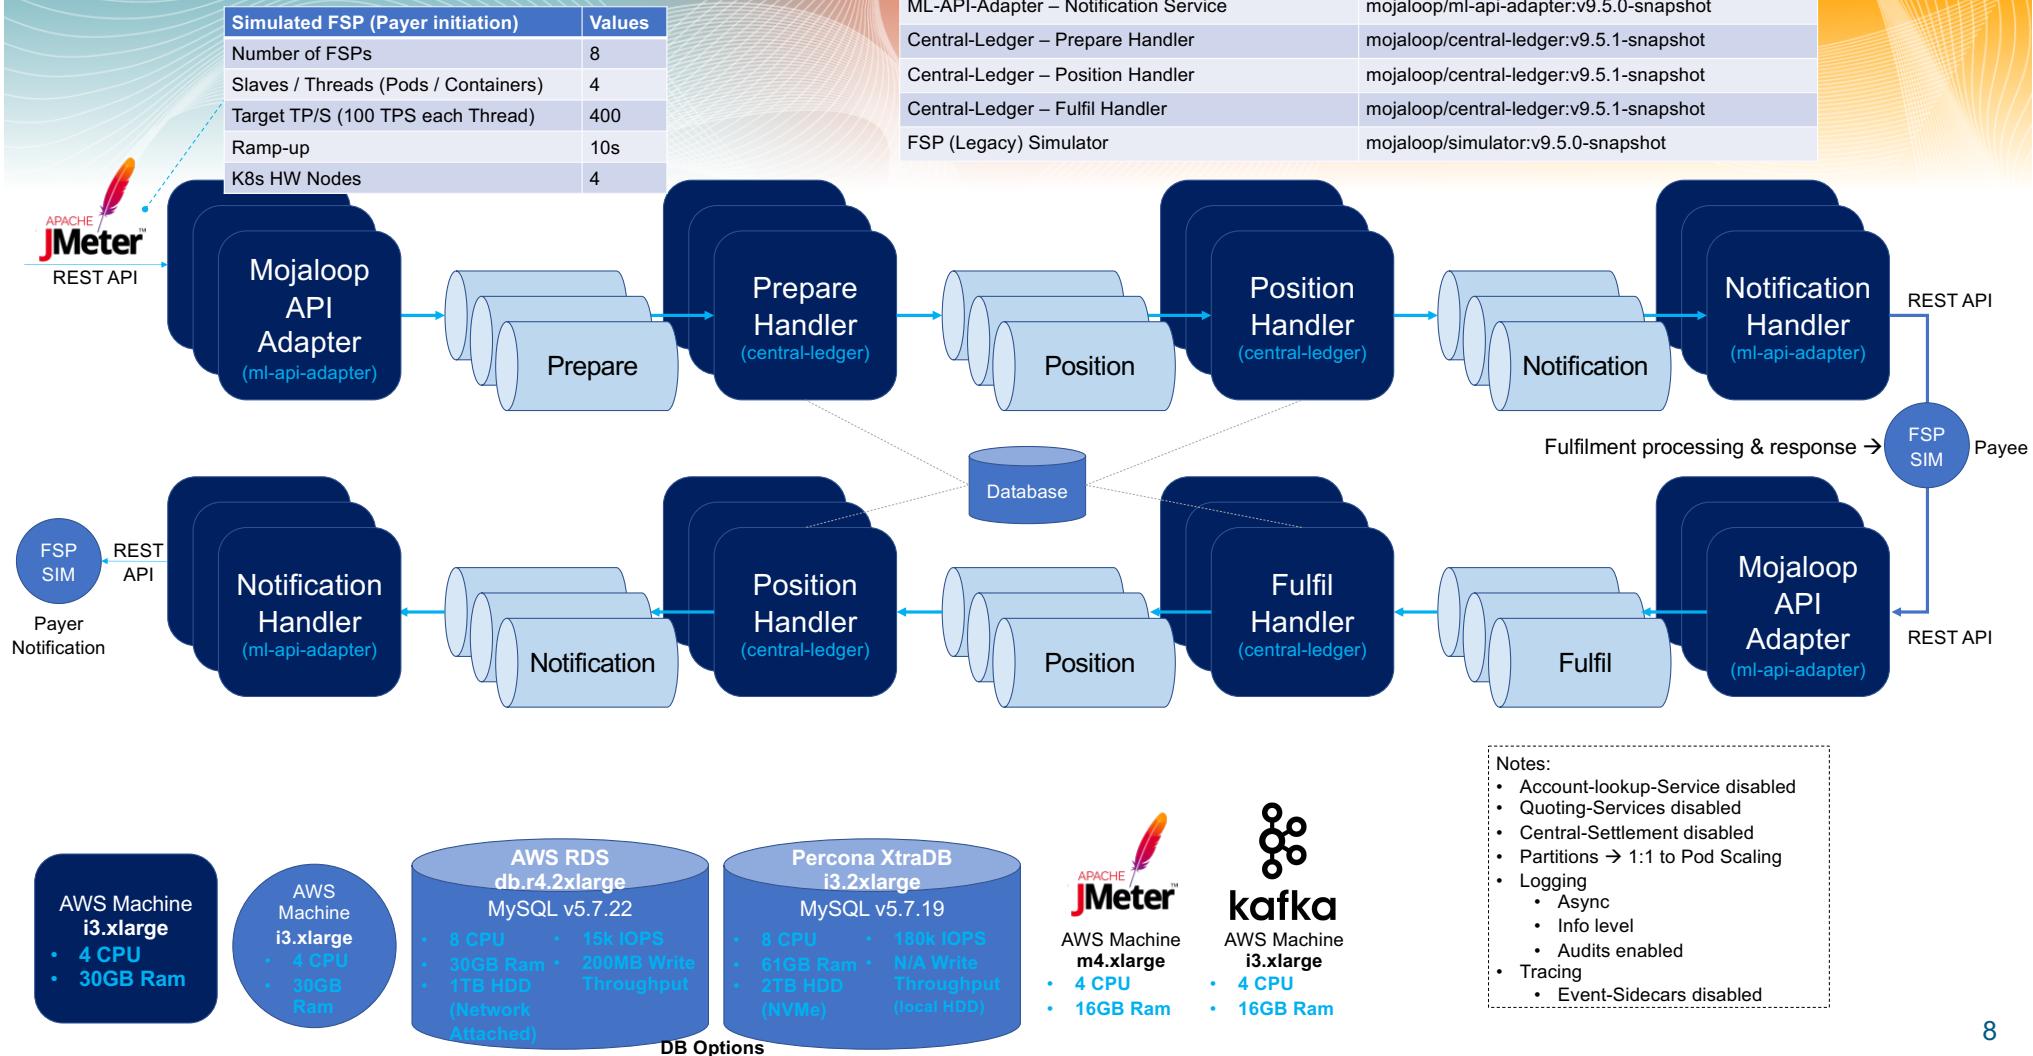
# Performance: Environment



## Environment

- Kubernetes Version: v1.11.6
- Rancher: v2.1.6
- Docker Version: 17.03.2-ce
- Kernel Version: 4.4.0-1052-aws
- Operating System: Ubuntu 16.04.4 LTS

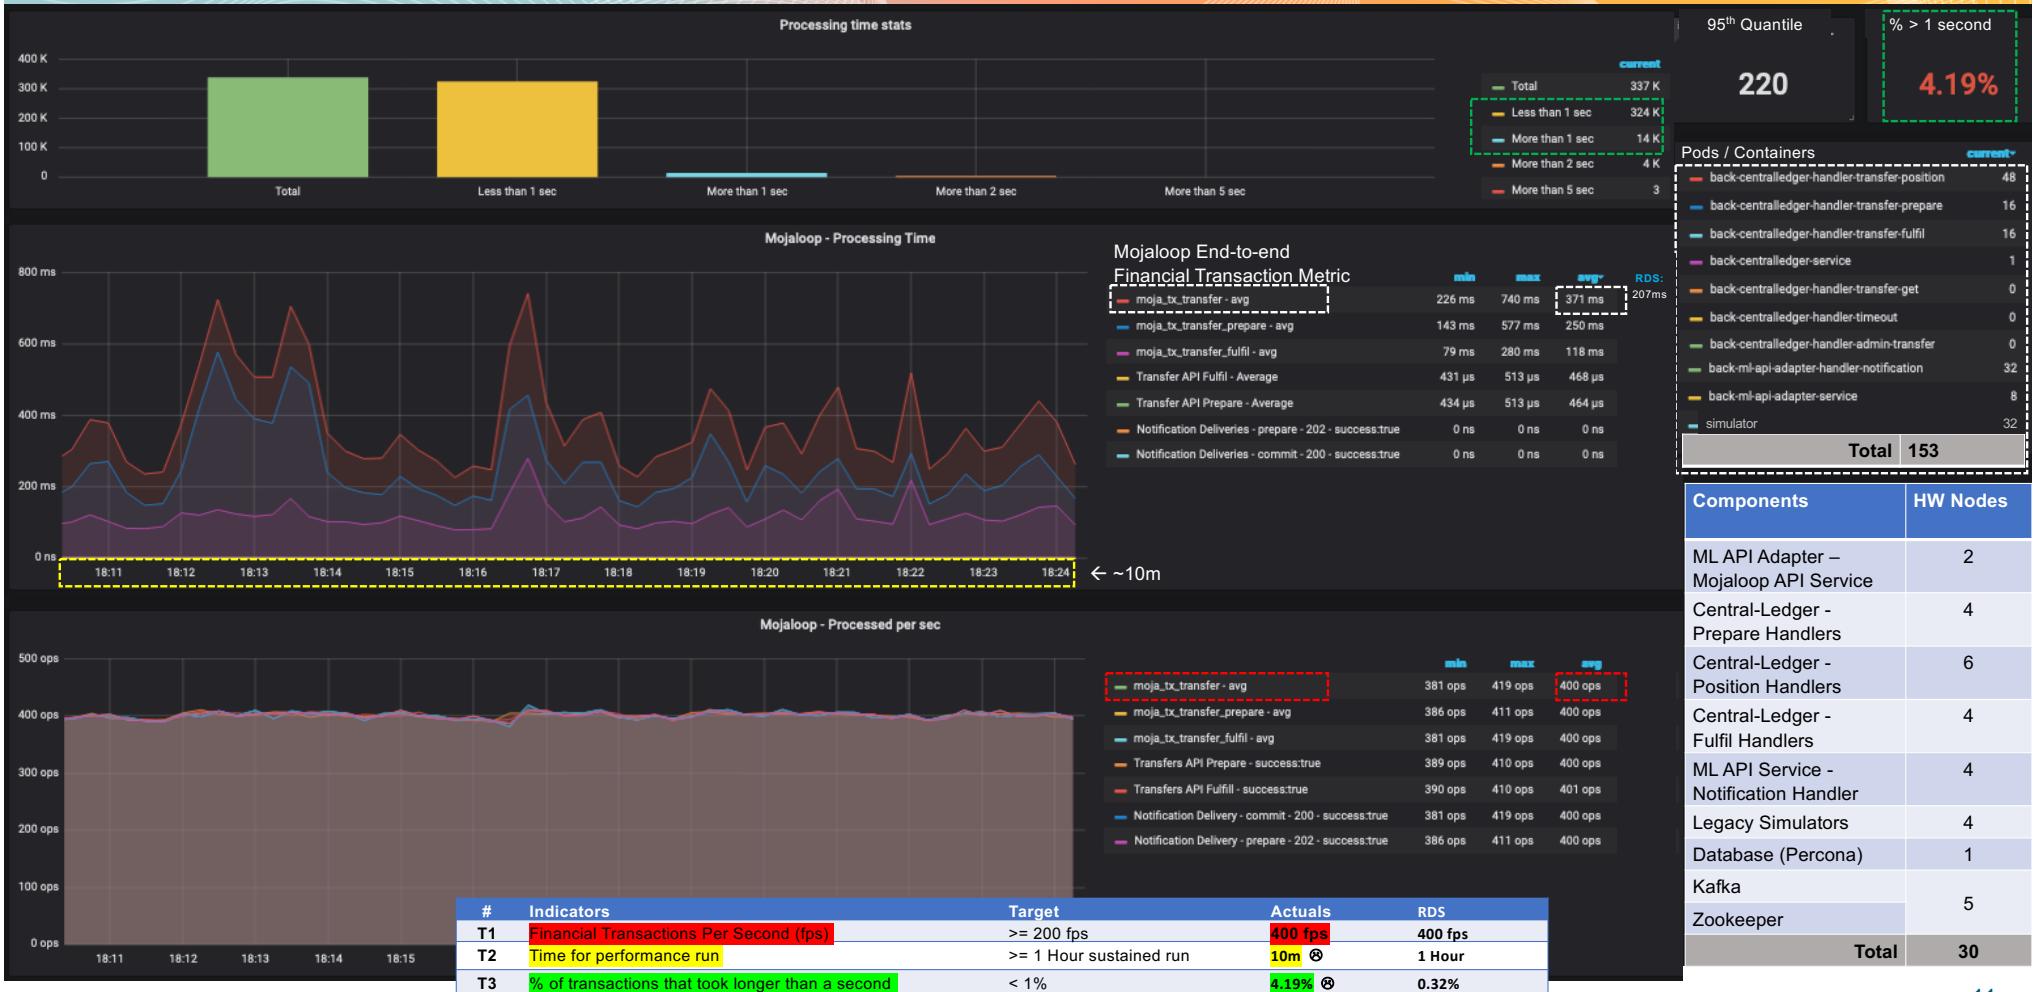
# Master Branch: Setup



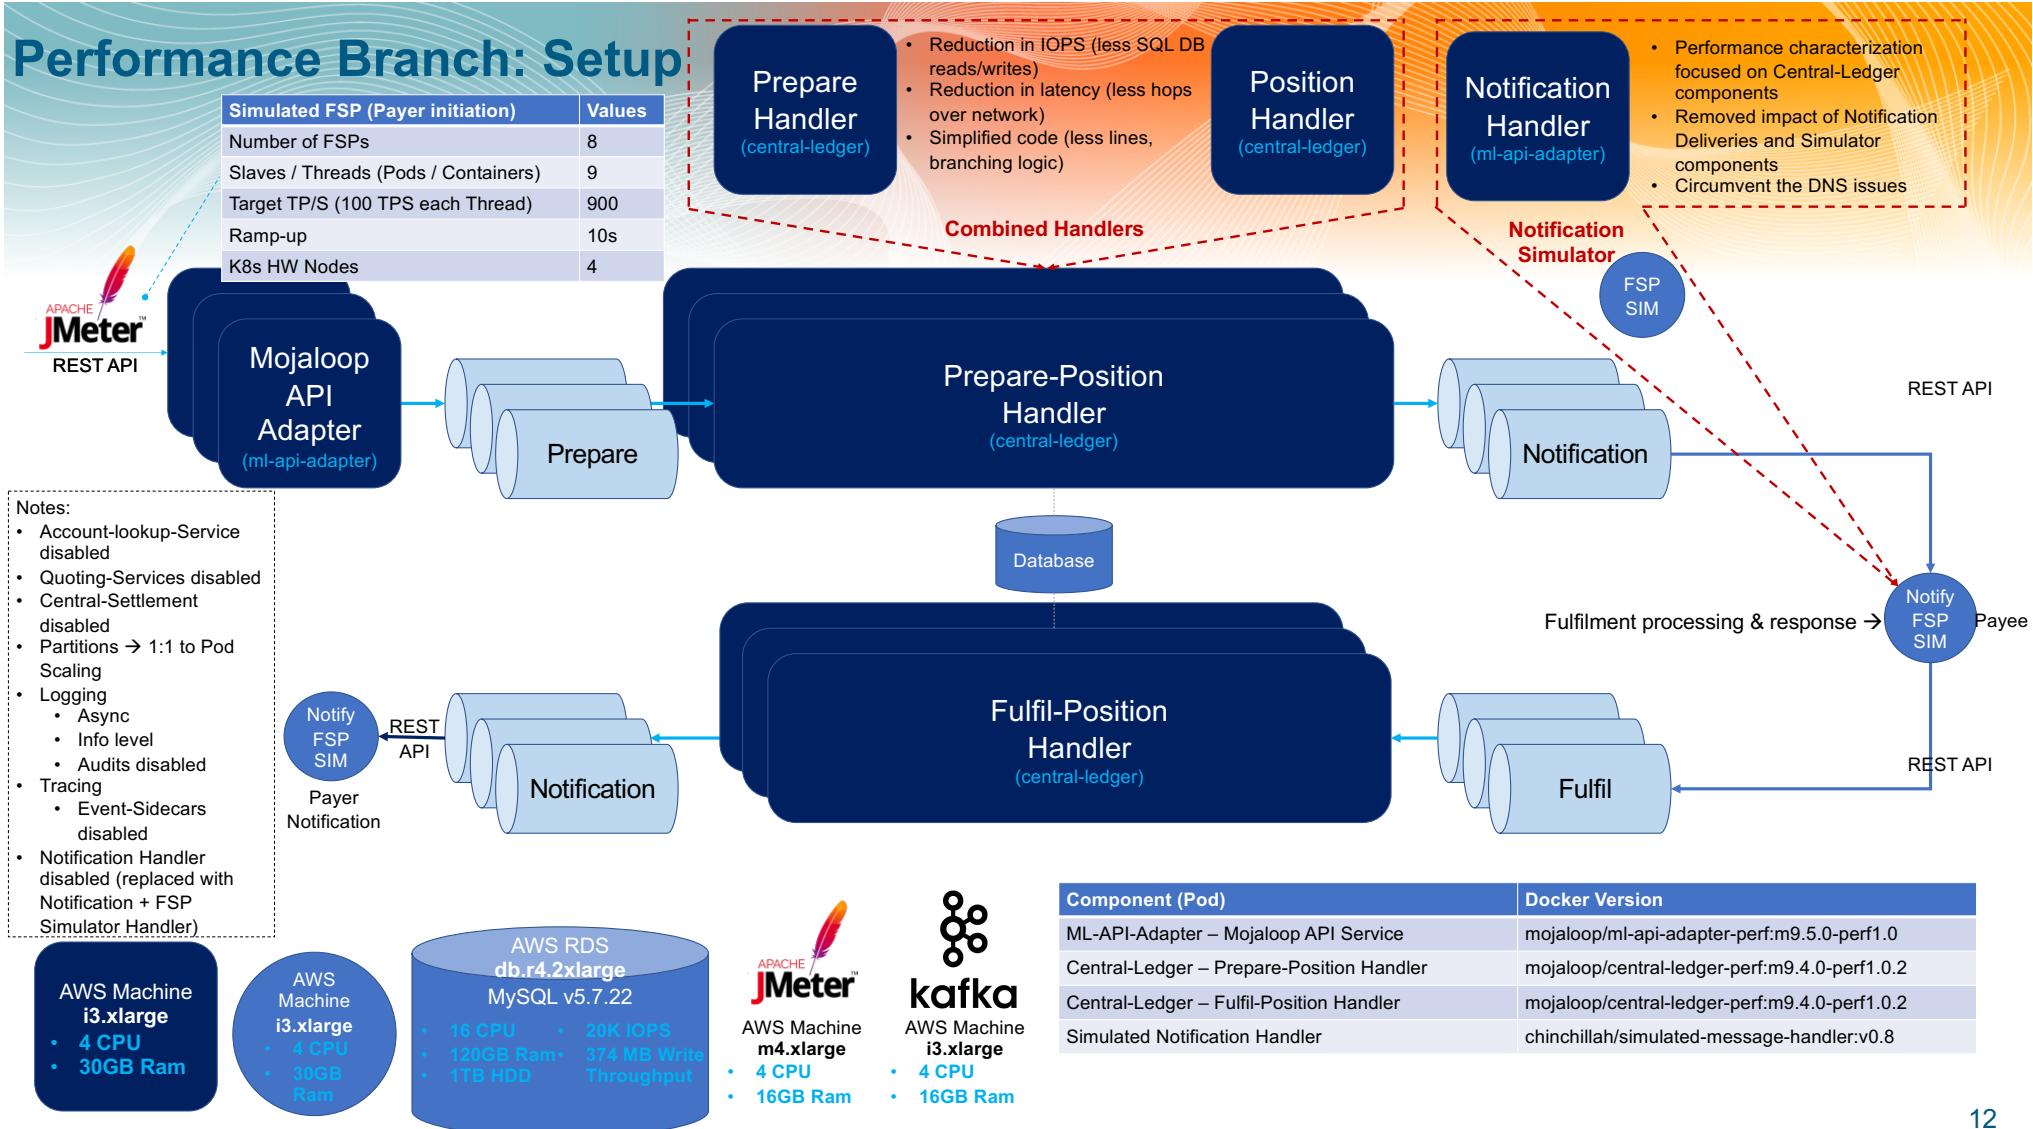
# Master Branch: Results @ 400 financial transaction per/second (fps) with RDS



# Master Branch: Results @ 400 financial transaction per/second (fps) self-hosted Percona XtraDB MySQL (unoptimized on-prem DB - early tentative result)



## Performance Branch: Setup



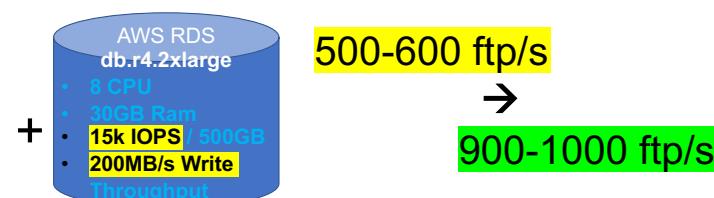
## Performance Branch: Comparison

1. Differences to Master (Changes to be reviewed, verified and productionized)
  - a) Combined Prepare + Position Handlers
  - b) Combined Fulfil + Position Handlers
  - c) Optimized SQL
    - a) Duplicate Checks
    - b) Position Changes
  - d) Full Caching for non-stateful data (participants)
  - e) Simulated Notification Handler (Direct-bypass to Fulfil topic from Notification Handler)

2. Improvements over Master with equivalent Database (db.r4.2xlarge @ 200MB/s Write Throughput)
  - a) Less Hardware Nodes: 30 → 17
  - b) Performance Improvement: ~20%+
  - c) Less IOPS (reduced SQL Queries)
  - d) Simpler code (improved maintainability)

### 3. Scaling Infrastructure on Performance Branch

| Components                                 | HW Nodes<br>(i3.xlarge) | Pods /<br>Container<br>s |
|--------------------------------------------|-------------------------|--------------------------|
| ML API Adapter – Mojaloop API Service      | 3                       | 20                       |
| Central-Ledger – Prepare-Position Handlers | 4                       | 34                       |
| Central-Ledger - Fulfil-Position Handlers  | 4                       | 34                       |
| Simulated Notification Handler             | 4                       | 20                       |
| Database (RDS)                             | 1                       | N/A                      |
| Kafka                                      |                         | 5                        |
| Zookeeper                                  | 5                       | 3                        |
| <b>Total</b>                               | <b>17</b>               | <b>125</b>               |

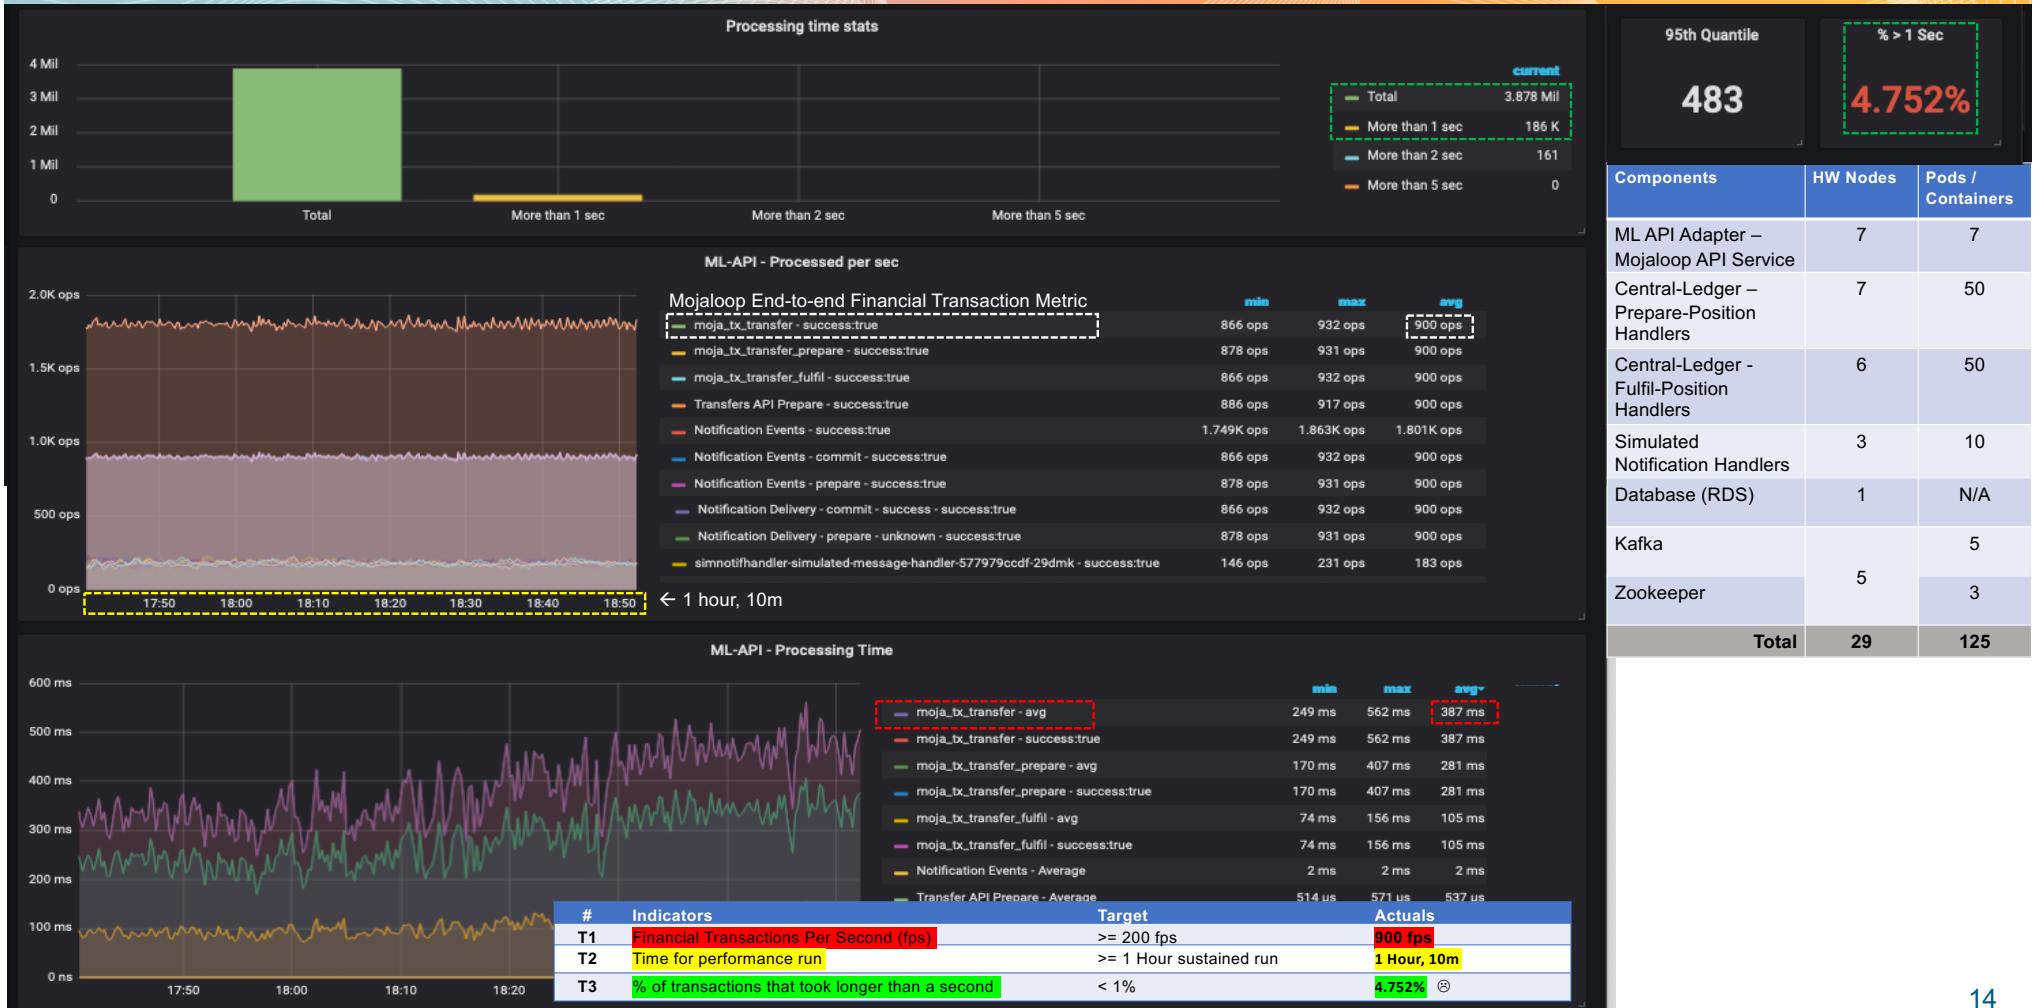


Notes:

- The hardware footprint has NOT been optimized for the Performance Branch

| Components                                 | HW Nodes<br>(i3.xlarge) | Pods /<br>Container<br>s |
|--------------------------------------------|-------------------------|--------------------------|
| ML API Adapter – Mojaloop API Service      | 7                       | 7                        |
| Central-Ledger – Prepare-Position Handlers | 7                       | 50                       |
| Central-Ledger - Fulfil-Position Handlers  | 6                       | 50                       |
| Simulated Notification Handler             | 3                       | 10                       |
| Database (RDS)                             | 1                       | N/A                      |
| Kafka                                      |                         | 5                        |
| Zookeeper                                  | 5                       | 3                        |
| <b>Total</b>                               | <b>29</b>               | <b>125</b>               |

# Performance Branch: Results @ financial transaction per/second (fps)



# Unresolved Issues

## 1. Notification Delivery

- a) DNS Issues due to issue on underlying Kubernetes OSS Environment (**i.e. not code related**) → temporarily resolved by going directly to Simulator's Service Cluster-IP address



## 2. Optimize On-prem MySQL Percona Xtra DB Cluster Deployment:

- a) Optimize DB configuration for underlying hardware:
  - a) Queue Cache Size
  - b) Utilize Available Memory
  - c) etc.
- b) Optimize infrastructure (e.g. AWS)
  - a) Ensure storage has guaranteed IOPS provisioned
- c) Resolve scaling issues
  - a) Refactor Transaction Logic to cater for optimistic locking (**note: not an issue on Azure/AWS managed DBs**)
  - b) Consider using alternative pessimistic distributed locking (e.g. Redlock) system in conjunction

## 3. Cloud Provider (AWS) Shared Infrastructure

- a) Provision dedicated instances for key systems (e.g. Kafka, Zookeeper, Database, etc)

# Roadmap - Future Enhancements

1. Optimize Run-time Environment
  - a) Running Kubernetes Cluster on dedicated hardware
  - b) Identify & Resolve DNS Issue
2. Optimize On-prem MySQL Percona Xtra DB Cluster Deployment:
  - a) Optimize DB configuration for underlying hardware:
    - a) Queue Cache Size
    - b) Available Memory
  - b) Optimize infrastructure (e.g. AWS)
    - a) Ensure storage has guaranteed IOPS provisioned
  - c) Resolve scaling issues
    - a) Refactor Transaction Logic to cater for optimistic locking
    - b) Consider using alternative pessimistic distributed locking (e.g. Redlock) system in conjunction
3. Optimize On-prem Kafka Deployment:
  - a) Separate Kafka / Zookeeper clusters
  - b) Ensure dedicated hardware
4. Optimize Notification Handlers / Simulators
  - a) Optimize HTTP Connections (keep-alive, connection-pool, etc)
5. Optimized SQL
  1. Duplicate Inserts
  2. Position Change
6. Caching
  - a) Productionize Participant Limits (central-ledger)
7. Metrics – Instrument End-to-end Financial metrics
  - a) Account-lookup
  - b) Quoting-service
  - c) Settlement-service
  - d) Etc
8. Optimize Logging on remaining components
  - a) Account-lookup
  - b) Quoting-service
  - c) Settlement-service
  - d) Etc
9. Timeout Handler causes an intermittent dead-lock during high-loads
  - a) Ensure no table scan locks (dirty reads)
10. Review & Productionize Combined Handler PoC

# Roadmap - Future Characterizations

## 1. Scalability

- a) Horizontal Scaling of Mojaloop
- b) FSP ← impact of increasing # of FSPs

## 2. Other Mojaloop Services/Flows

- a) Account Lookup
- b) Quoting
- c) Settlements
- d) Transactions-requests
- e) Unhappy Paths

## 3. Event-Framework

- a) Event-Sidecar
- b) Event-Stream-Processor
- c) ElasticSearch (Event / Log ingestions)
- d) APM (Trace flow reporting)

## 4. API Gateway

## 5. Security

- a) JWS
- b) TLS

## 6. Simulators

- a) SDK-Schema-Adapter

## 7. Proof of Concepts

- a) Compare against baselines: master and performance branches

# mojaloop

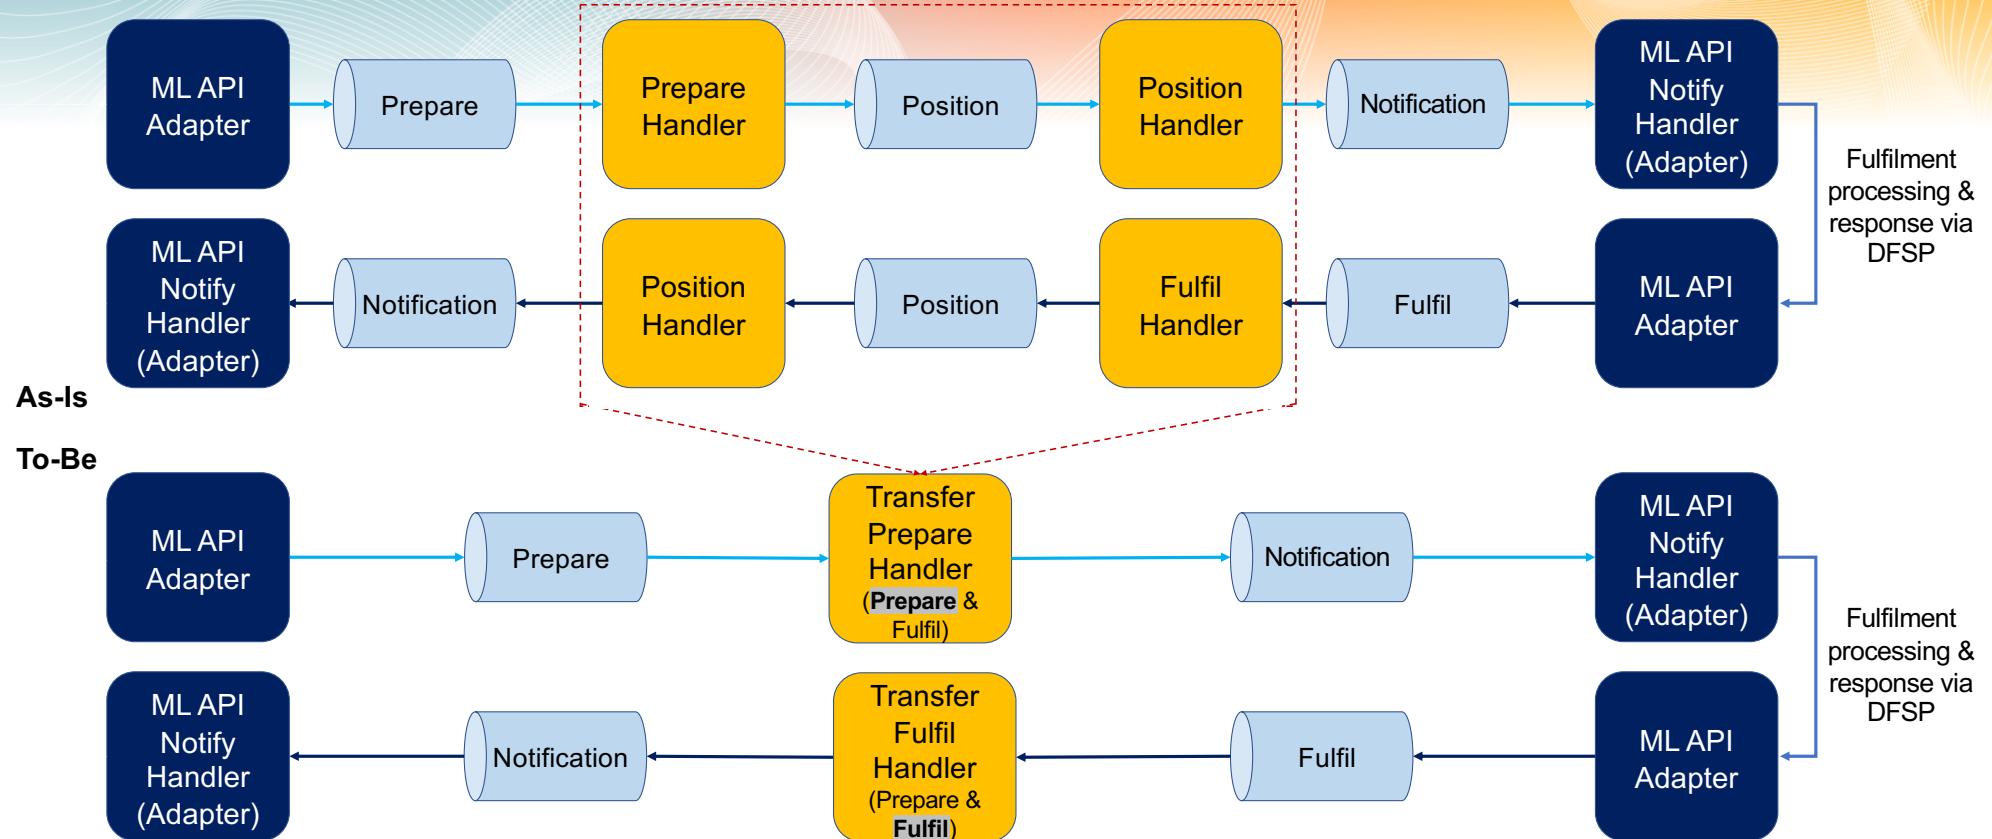
## Questions?

mojaloop

# Performance

APPENDIX

# Performance Branch: Combined Central Handlers



- Reduction in IOPS (less SQL DB reads/writes)
- Reduction in latency (less hops over network)
- Simplified code (less lines, branching logic)

# In-Memory Bulk Processing

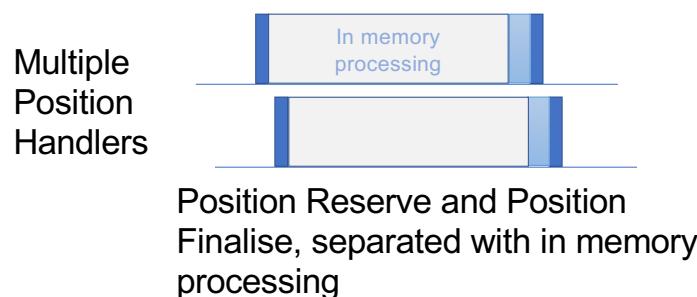
Objective: Compete historic state record at time of rule execution, while never breeching liquidity checks

## 1. Position Handler Turnstile – single instance, bulk transfer set for processing



1. DFSP's Position is held in READ/UPDATE lock while processing transfer set read from topic
2. Means there can only be one active handler per DFSP

## 2. Position Reserved Handler – Reserved Position, multi-instances



1. Position/Liquidity is reserved for total of transfer set
2. After in memory processing:
  - a. Finalised transfer states are INSERTed, and
  - b. DFSP Liquidity is finalised with UPDATE
3. Works in mixed or per DFSP processing mode
4. Multiple active handlers