



Technical Overview

Mojaloop PI10 Pre-Workshop

October 2020

Miguel de Barros
miguel.debarros@modusbox.com
v9.5.x



Mojaloop Links

Website: <http://mojaloop.io>

Docs: <http://mojaloop.io/documentation>

Slack: <http://mojaloop.io/slack> ([self-invite](#))

Code: <http://github.com/mojaloop>

Gitter: <https://gitter.im/mojaloop/mojaloop>

October 2020

7

Mojaloop

Technical Overview

1. What is Mojaloop?
2. High-level Overview, Static Demo & Architecture
3. Component Architecture
4. Global foot-print & Roll-out

Mojaloop

Technical Overview

1. What is Mojaloop?
2. High-level Overview, Static Demo & Architecture
3. Component Architecture
4. Global foot-print & Roll-out

What is Mojaloop?

- Open Loop System
- Real-time, Irrevocable, Push-only
- DFSP Governed, Same-day Settlements
- Shared investment in Fraud detection

The logo consists of a white lowercase 'o' centered within three concentric dashed circles. Below the graphic, the word 'mojaloop' is written in a large, lowercase, sans-serif font.

mojaloop

October 2020

Mojaloop

Technical Overview

1. What is Mojaloop?
2. High-level Overview, Static Demo & Architecture
3. Component Architecture
4. Global foot-print & Roll-out

Definition of general terms

1. Hub

- The general term for an operating Mojaloop Switch in the “wild”
- This includes Mojaloop and all supporting services/components/rules required to operate a fully functional payment switch
- Adheres to regulatory requirements when/where applicable

2. Scheme

- The rules around an operating Mojaloop Hub

- Limits
- Reporting
- Regulations
- Onboarding process
- etc

3. (D)FSP

- (Digital) Financial Service Provider

4. Participants

- General term for an FSP that utilizes Mojaloop to exact payments

5. Party

- The customer that belongs to an FSP(s) which is represented by an identifier (e.g. MSISDN)

Definition of general terms

6. Cryptographic Receipt (<https://interledger.org/>)

- Cryptographic receipt used to validate payments which comprised of two parts:
 - Fulfilment ← cryptographically generated from payment information by the Payee FSP
 - Condition ← hash of the Fulfilment (used to validate Transfers when Fulfilment is shared)

7. Participant Position

- A participant (FSP) position is a net position of all activity
- Applies to a switch account for a scheme participant in a specific currency

8. Net Debit Cap Limit

- The limit applied to a participant's position, i.e. a participant is not allowed to trade in the system if their position exceeds their Net Debit Cap Limit.

Stages of a Mojaloop P2P Transaction

1. Discovery

- a. Resolve link between a Party (Customer) and a Participant (FSP)
- b. Lookup Party details (name, account, etc)

2. Agreement

- a. Participants
- b. Transfer Amount
- c. Fees
- d. Generate Cryptographic Receipt (used for validation on Transfer phase)
 - a. Fulfilment → crypto receipt generated from payment information
 - b. Condition → hash of the Fulfilment

3. Transfer

- a. Request (Prepare)
 - i. Sub-set of Quote information (e.g. participants, amount, etc)
 - ii. Condition
- b. Response (Fulfil)
 - i. Commit / Abort
 - ii. Fulfilment (Crypt-Receipt)

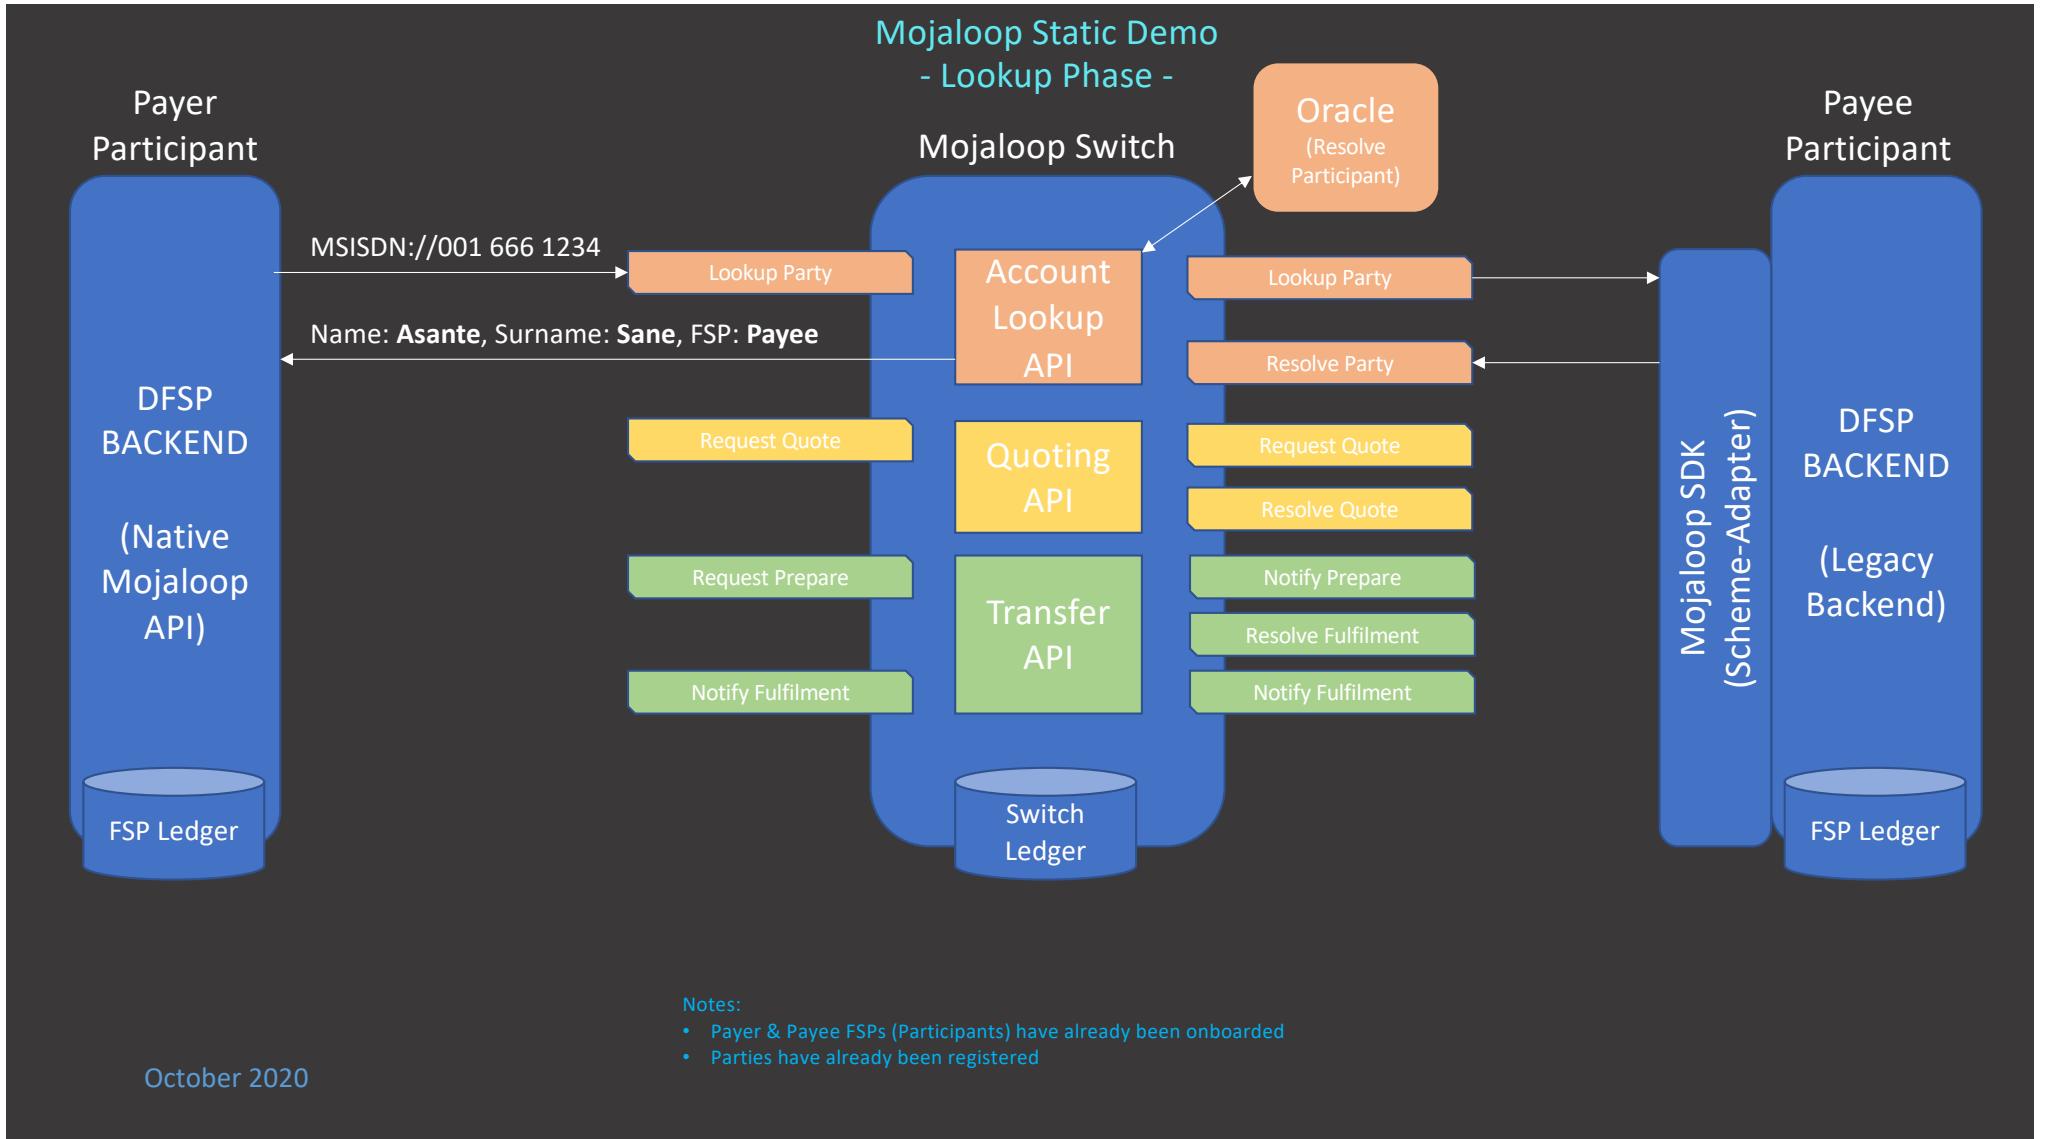
Static Demo: P2P Scenario

A customer ([Sending Party](#)) from an FSP ([Payer Participant](#)) wants to make a payment to a friend (*Asante Sane* – [Recipient Party](#)) that belongs to another FSP ([Payee Participant](#))

1. The [Payer Participant](#) does not know which FSP *Asante Sane* is a customer of. ←
[Payer Participant](#) can resolve this first via an [Account Lookup](#)
2. [Payer Participants](#) wants to agree/validate the [final transfer](#) (inc. costs, fees, etc) amount with its paying customer ([Sending Party](#)) ← [Payer Participant](#) can resolve this via a [Quote Request](#)
3. [Payer Participant](#) can exact this by a [Transfer Request](#)

Assumptions:

1. [Payer and Payee FSPs](#) are already members of the scheme
2. [Payer and Payee FSPs](#) have provided up-front collateral (and Position Limits have been set appropriately)



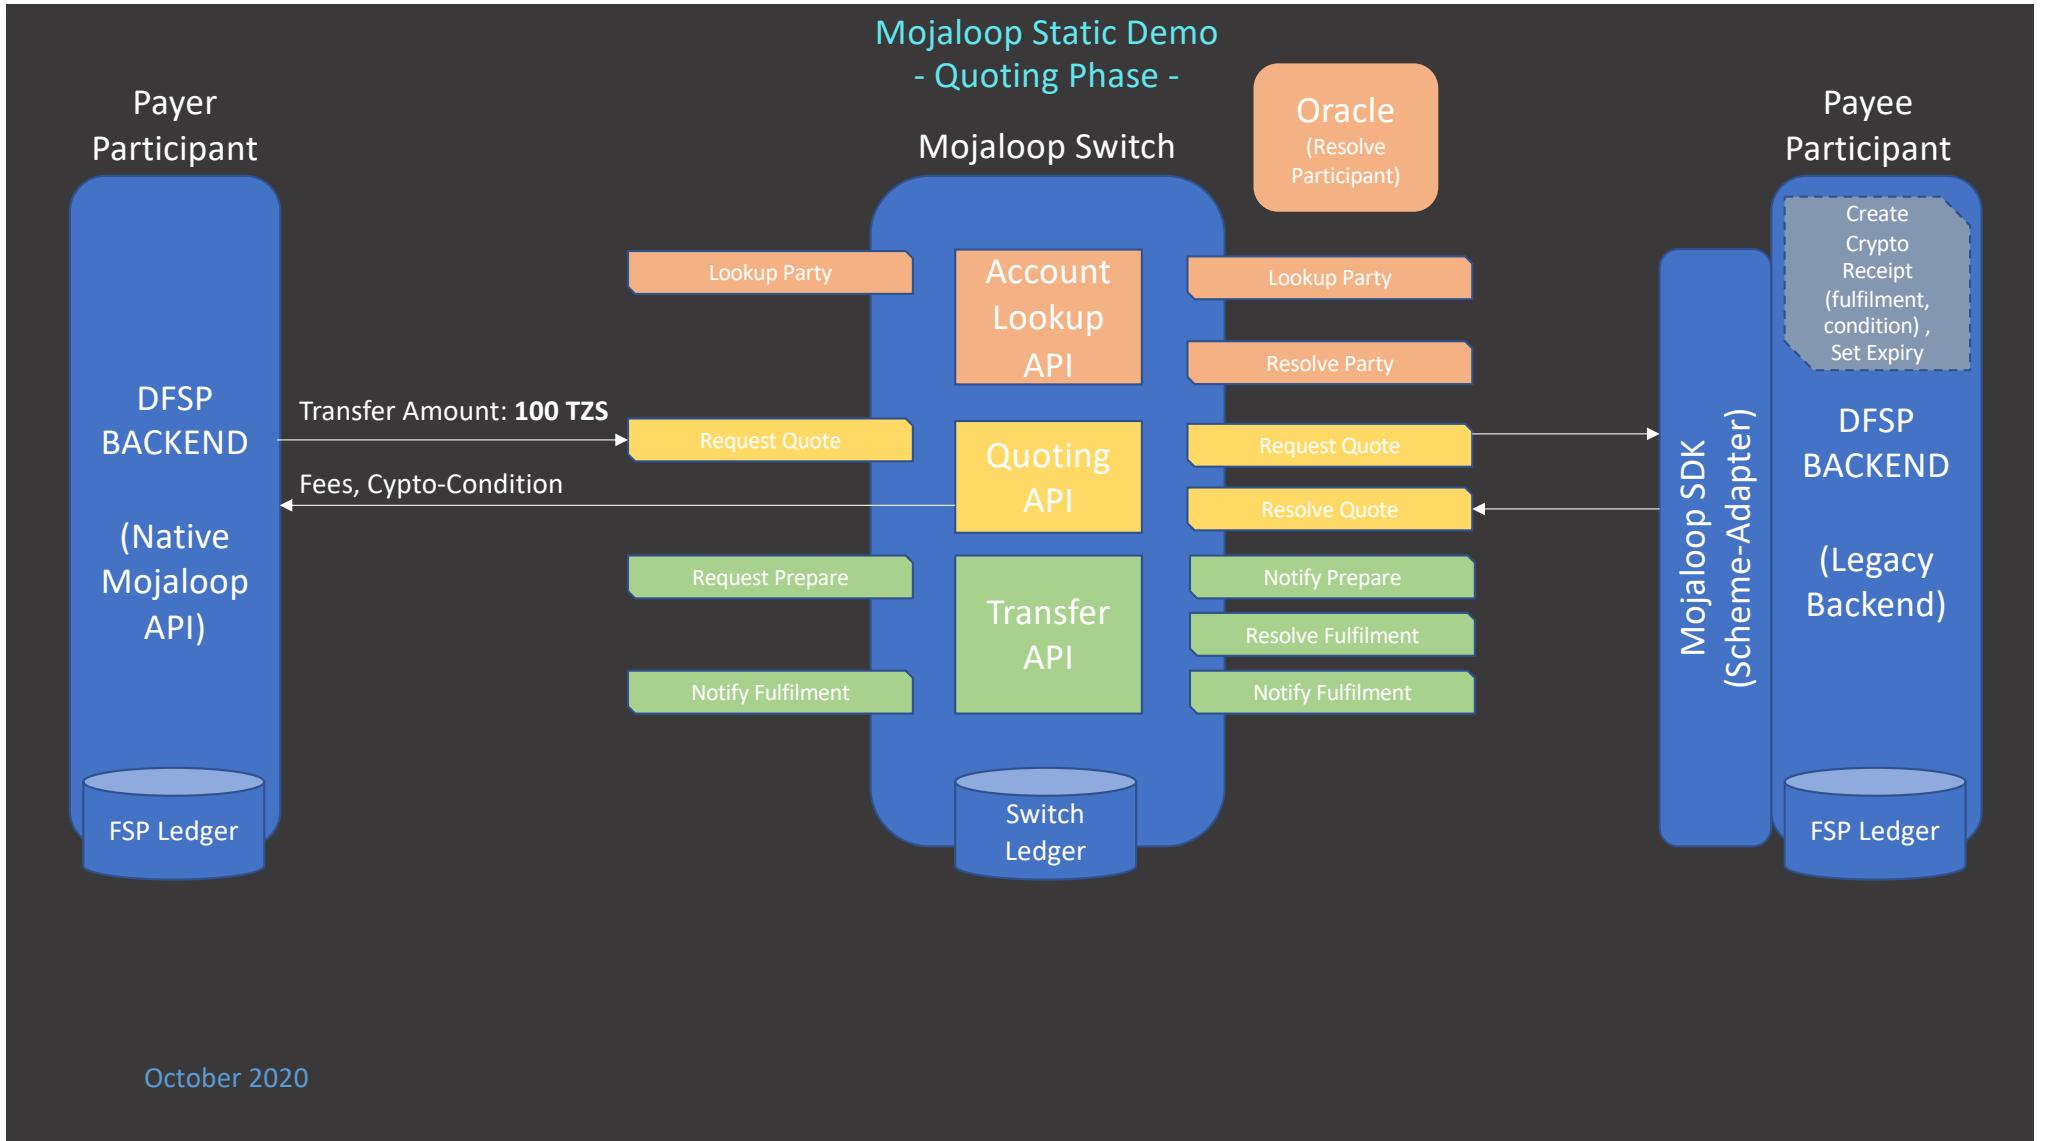
Static Demo: P2P Scenario

A customer ([Sending Party](#)) from an FSP ([Payer Participant](#)) wants to make a payment to a friend (*Asante Sane* – [Recipient Party](#)) that belongs to another FSP ([Payee Participant](#))

1. ~~The [Payer Participant](#) does not know which FSP *Asante Sane* is a customer of.~~ ←
[Payer Participant](#) can resolve this first via an [Account Lookup](#)
2. [Payer Participants](#) wants to agree/validate the [final transfer](#) (inc. costs, fees, etc) amount with its paying customer ([Sending Party](#)) ← [Payer Participant](#) can resolve this via a [Quote Request](#)
3. [Payer Participant](#) can exact this by a [Transfer Request](#)

Assumptions:

1. [Payer and Payee FSPs](#) are already members of the scheme
2. [Payer and Payee FSPs](#) have provided up-front collateral (and Position Limits have been set appropriately)



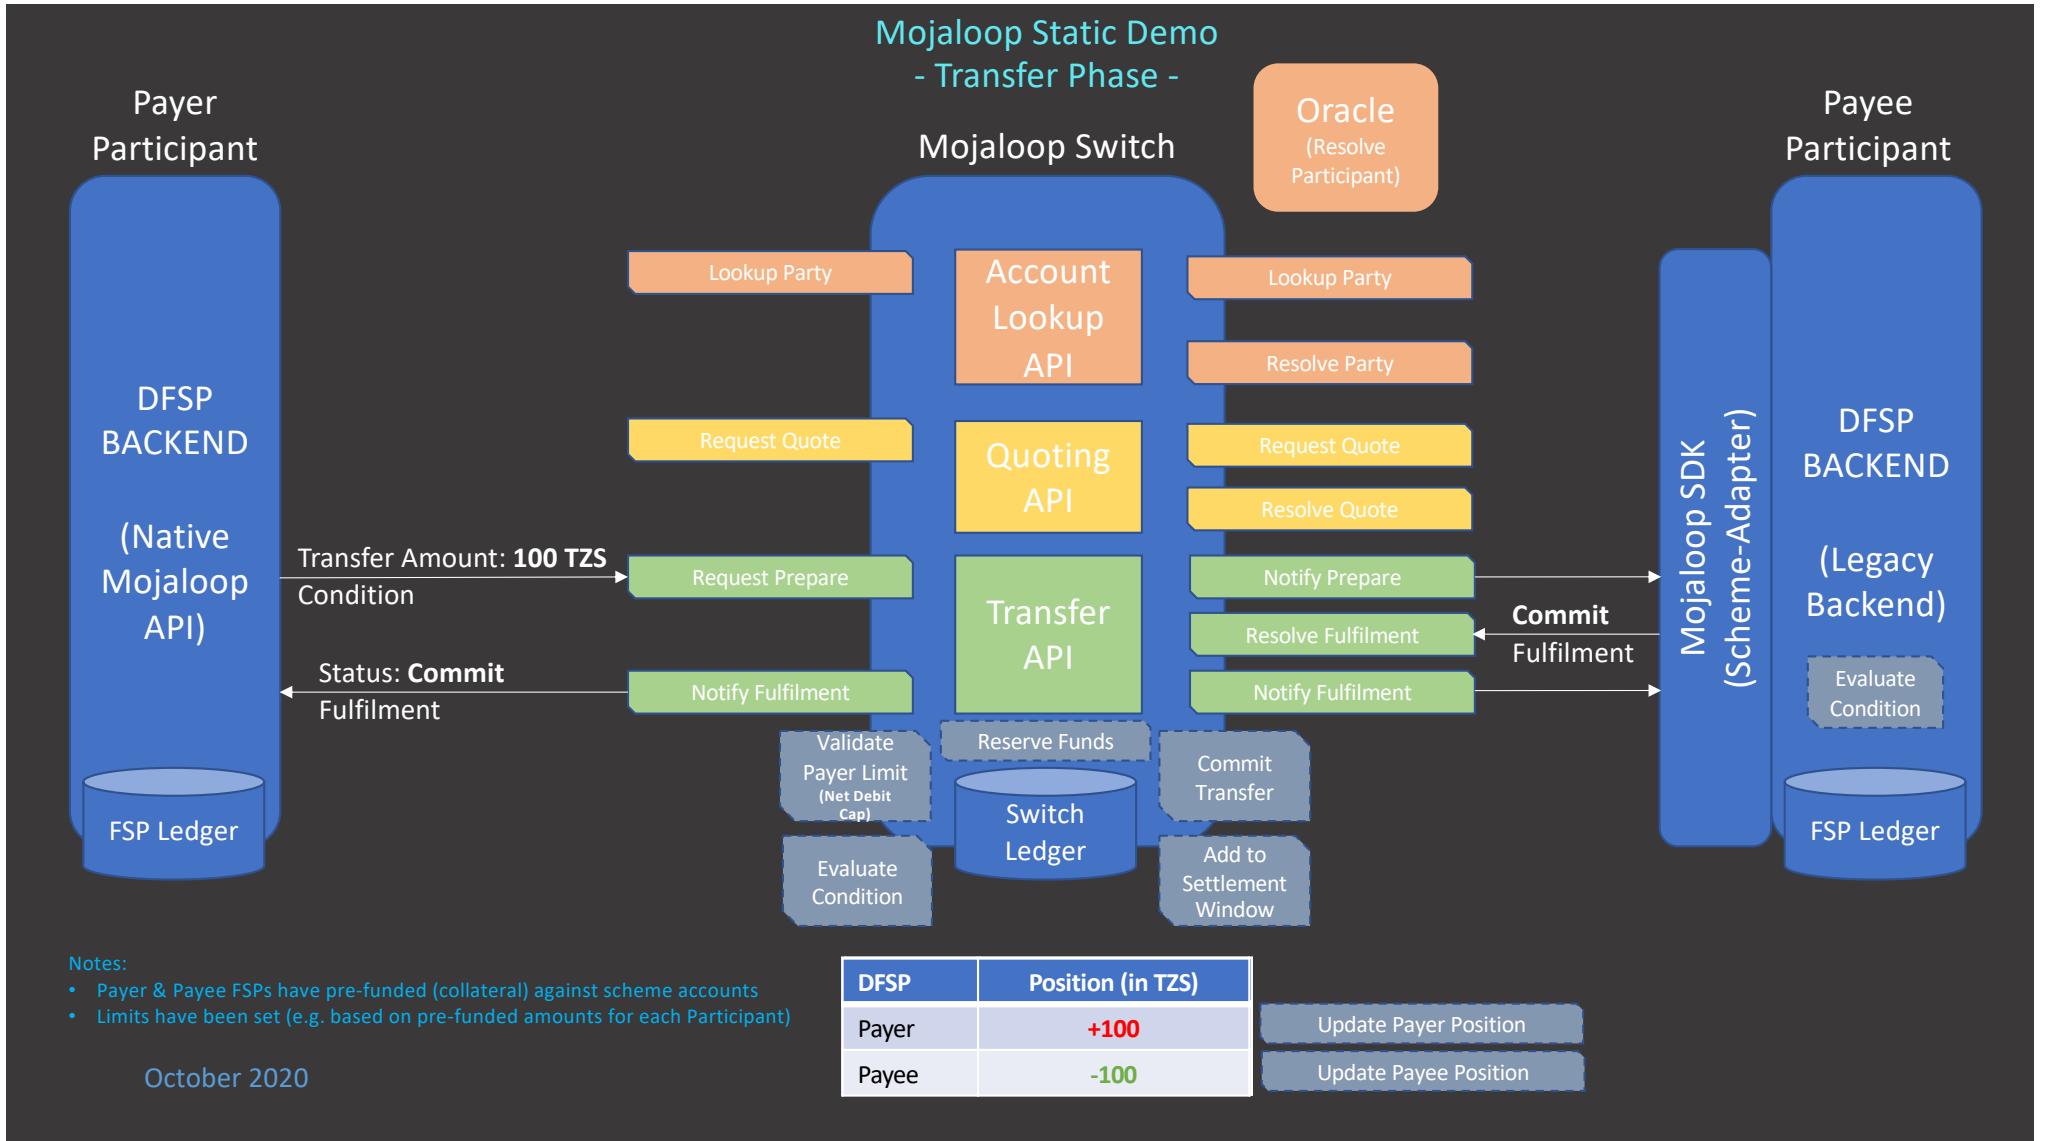
Static Demo: P2P Scenario

A customer ([Sending Party](#)) from an FSP ([Payer Participant](#)) wants to make a payment to a friend (*Asante Sane* – [Recipient Party](#)) that belongs to another FSP ([Payee Participant](#))

1. ~~The [Payer Participant](#) does not know which FSP *Asante Sane* is a customer of.~~ ←
[Payer Participant](#) can resolve this first via an [Account Lookup](#)
2. ~~Payer Participants wants to agree/validate the final transfer (inc. costs, fees, etc) amount with its paying customer ([Sending Party](#))~~ ← [Payer Participant](#) can resolve this via a [Quote Request](#)
3. [Payer Participant](#) can exact this by a [Transfer Request](#)

Assumptions:

1. [Payer and Payee FSPs are already members of the scheme](#)
2. [Payer and Payee FSPs have provided up-front collateral \(and Position Limits have been set appropriately \)](#)



Static Demo: P2P Scenario

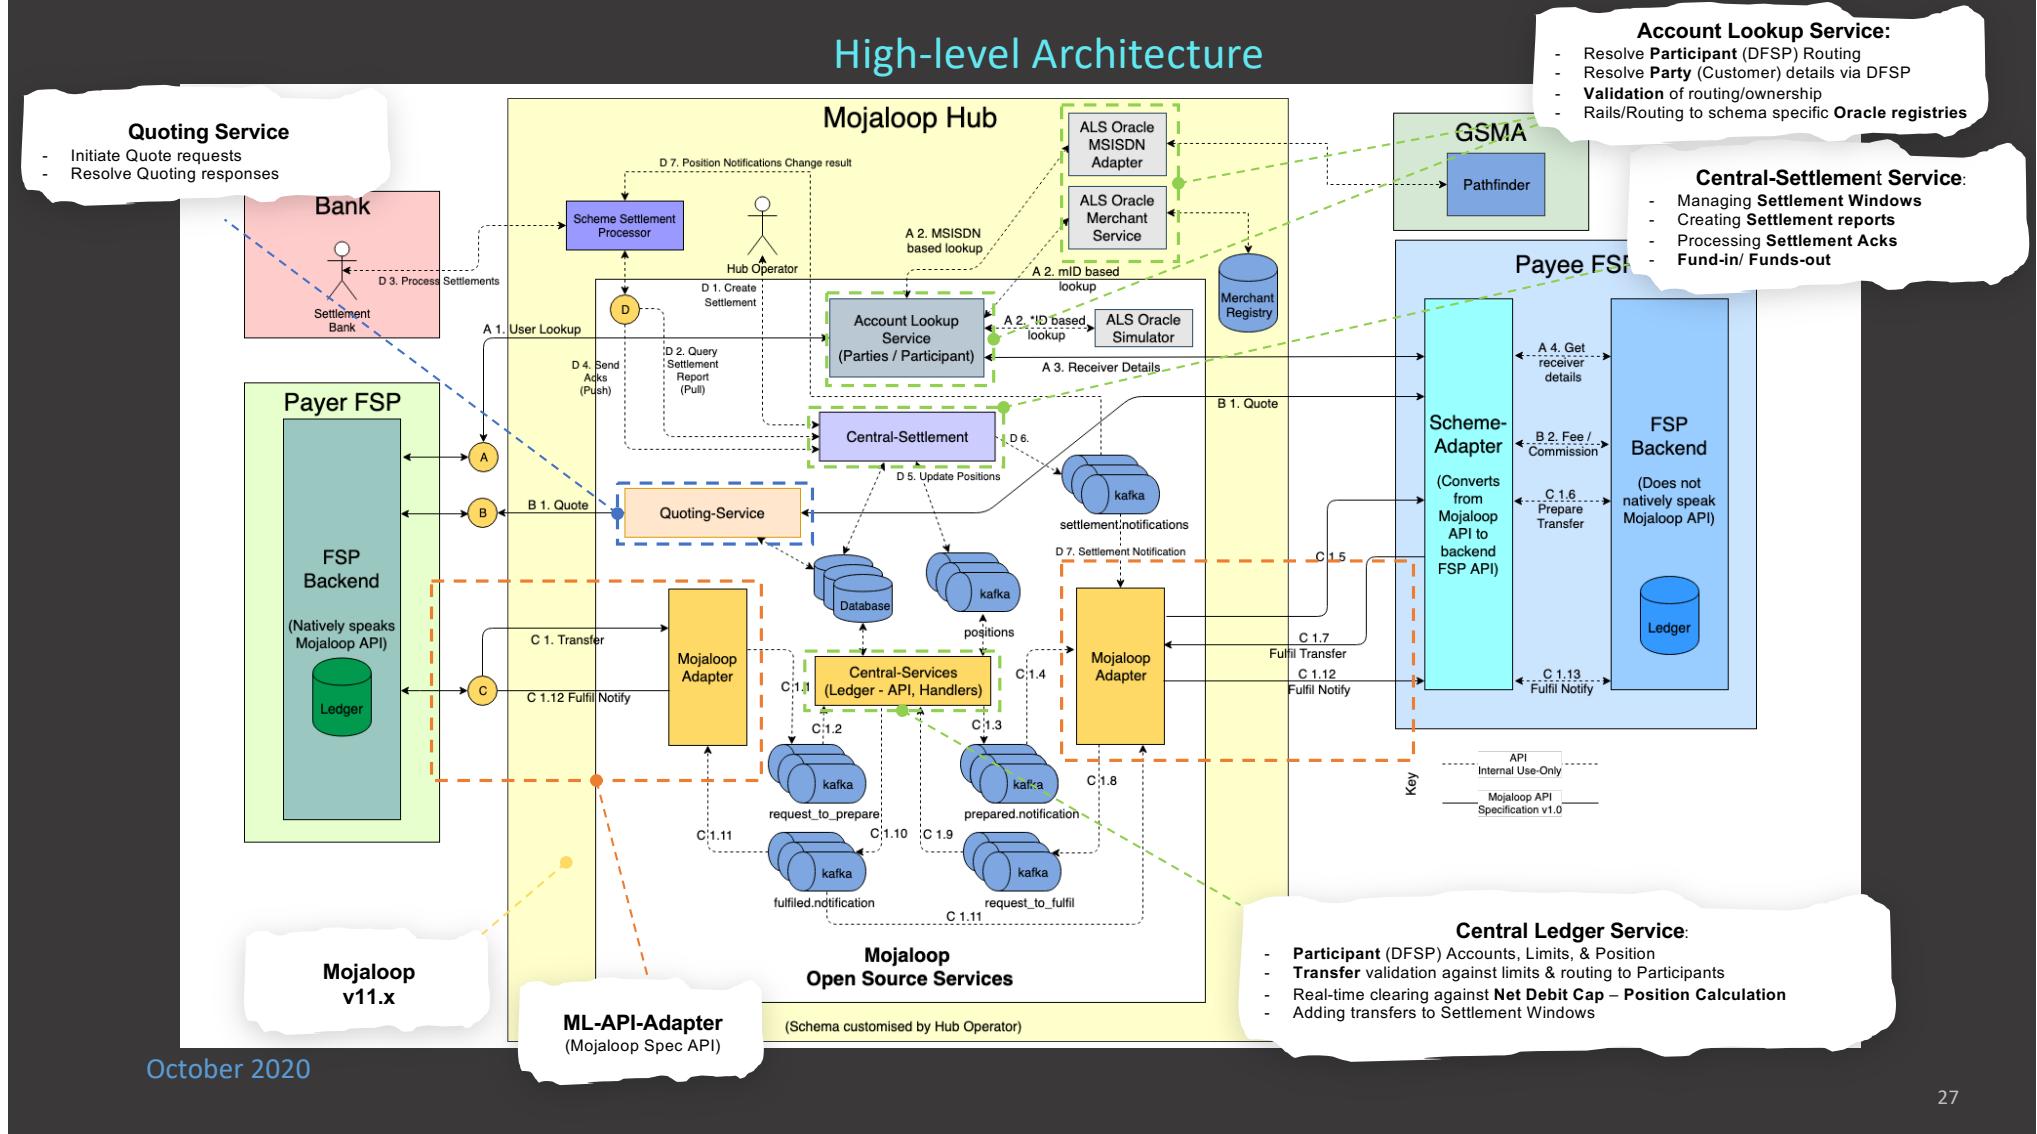
A customer ([Sending Party](#)) from an FSP ([Payer Participant](#)) wants to make a payment to a friend (*Asante Sane* – [Recipient Party](#)) that belongs to another FSP ([Payee Participant](#)) 

1. ~~The [Payer Participant](#) does not know which FSP *Asante Sane* is a customer of.~~ ←
[Payer Participant](#) can resolve this first via an [Account Lookup](#)
2. ~~Payer Participants wants to agree/validate the final transfer (inc. costs, fees, etc) amount with its paying customer ([Sending Party](#))~~ ← [Payer Participant](#) can resolve this via a [Quote Request](#)
3. ~~Payer Participant can exact this by a [Transfer Request](#)~~

Assumptions:

1. Payer and Payee FSPs are already members of the scheme
2. Payer and Payee FSPs have provided up-front collateral (and Position Limits have been set appropriately)

High-level Architecture



October 2020

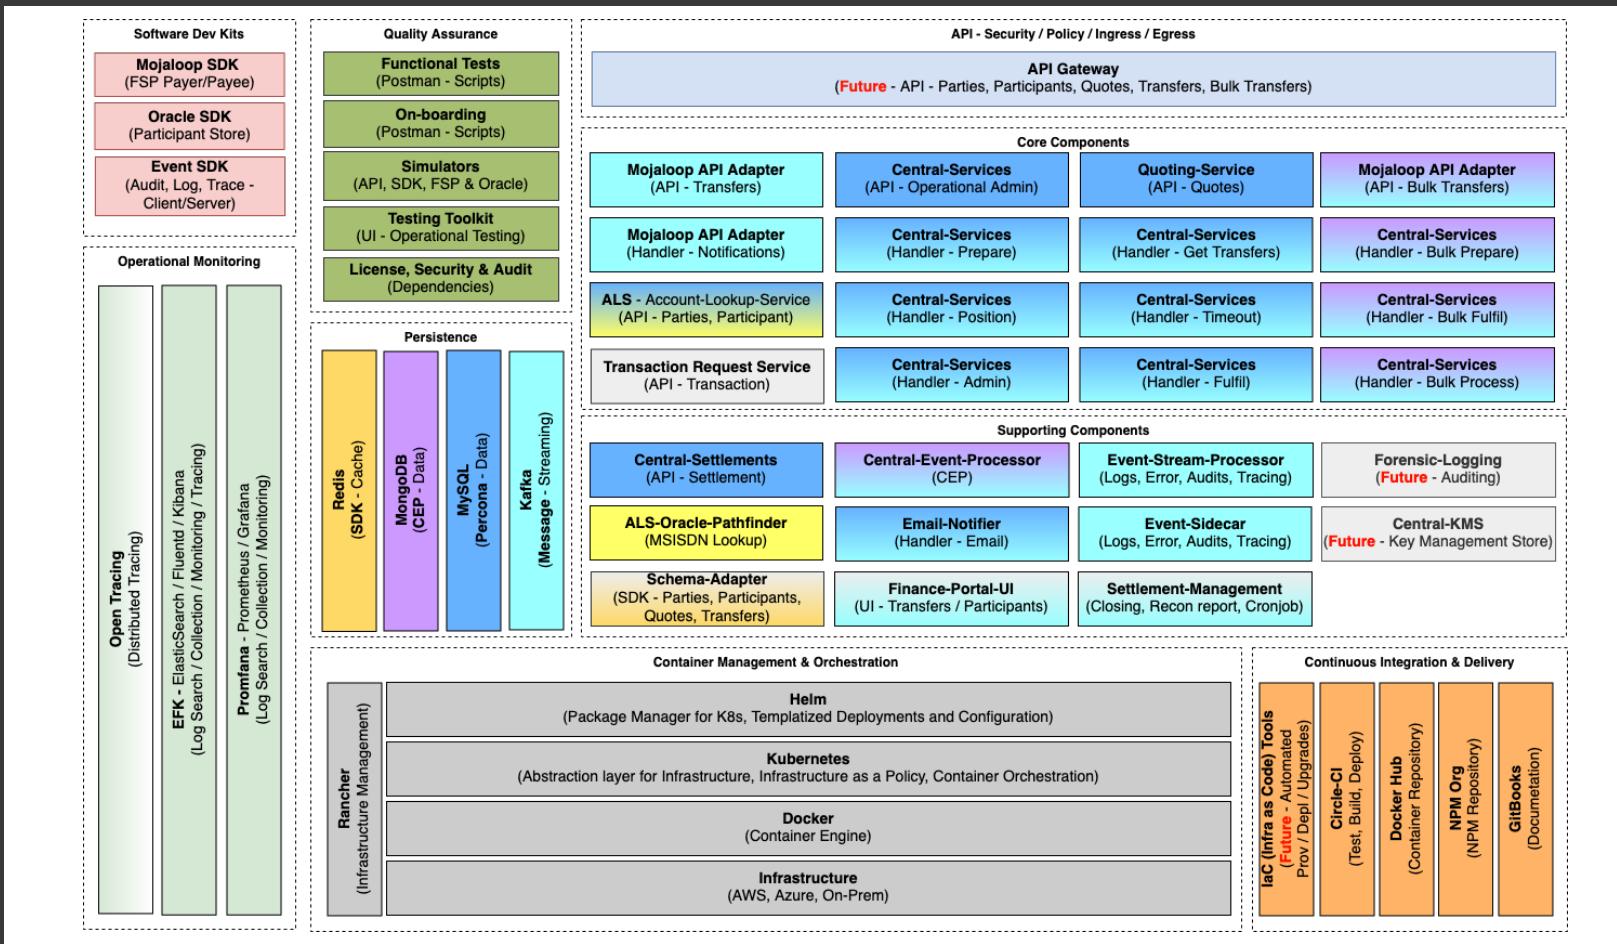
27

Mojaloop

Technical Overview

1. What is Mojaloop?
2. High-level overview, Static Demo & Architecture
3. [Component Architecture](#)
4. Global foot-print & Roll-out

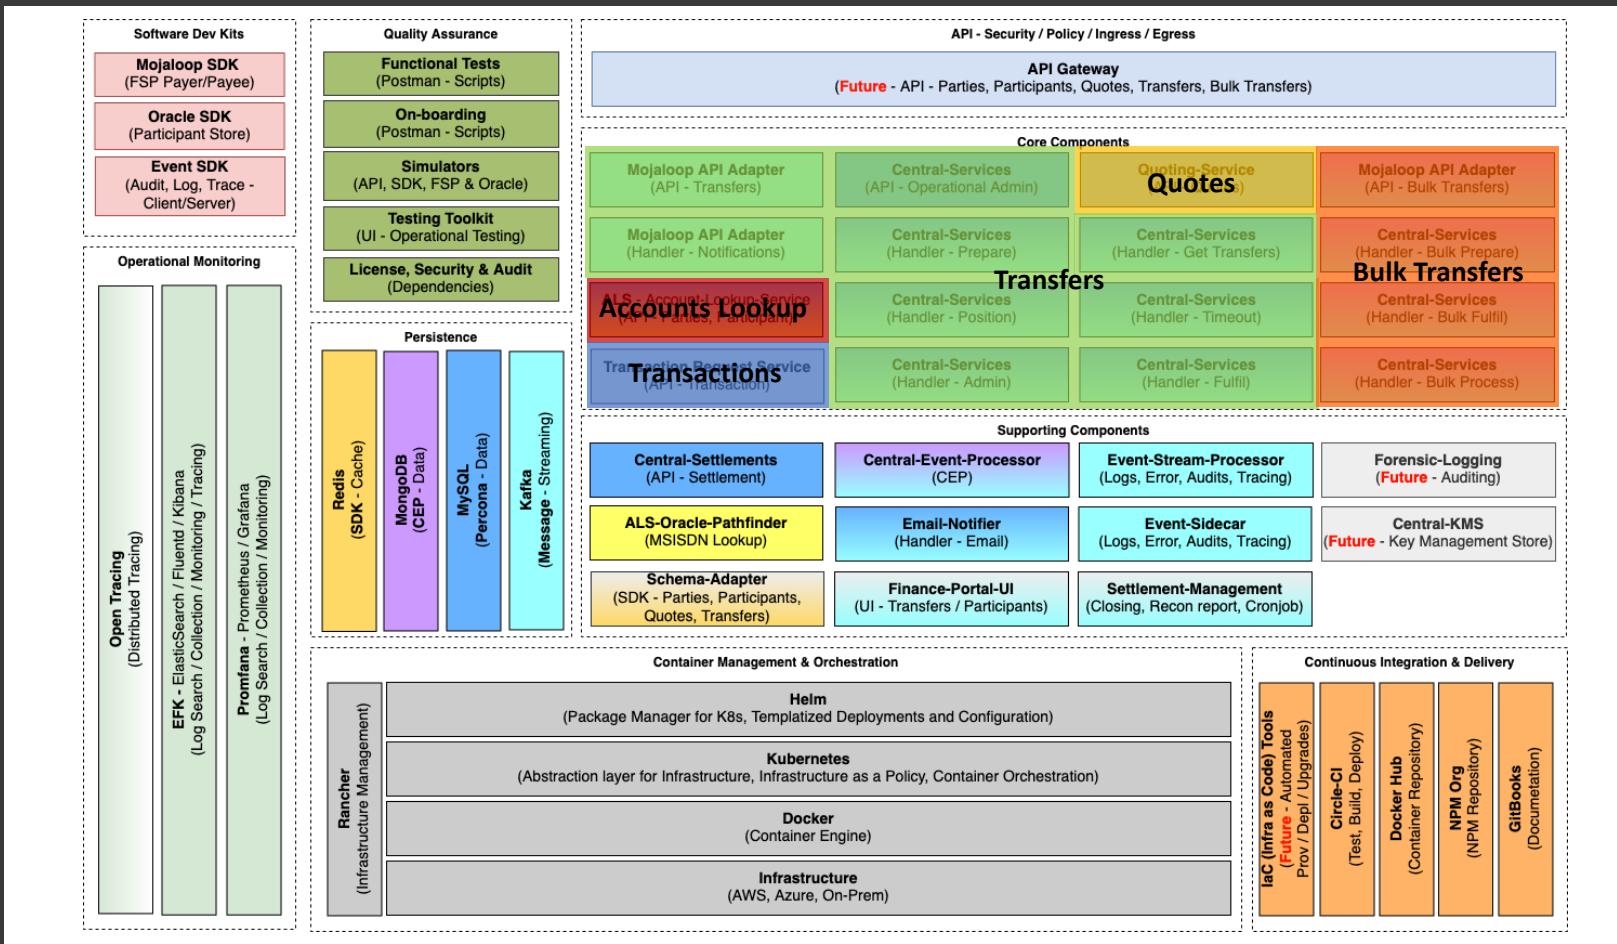
Component Architecture



October 2020

* Color-shading of core & supporting components match dependent persistent stores.

Component Architecture

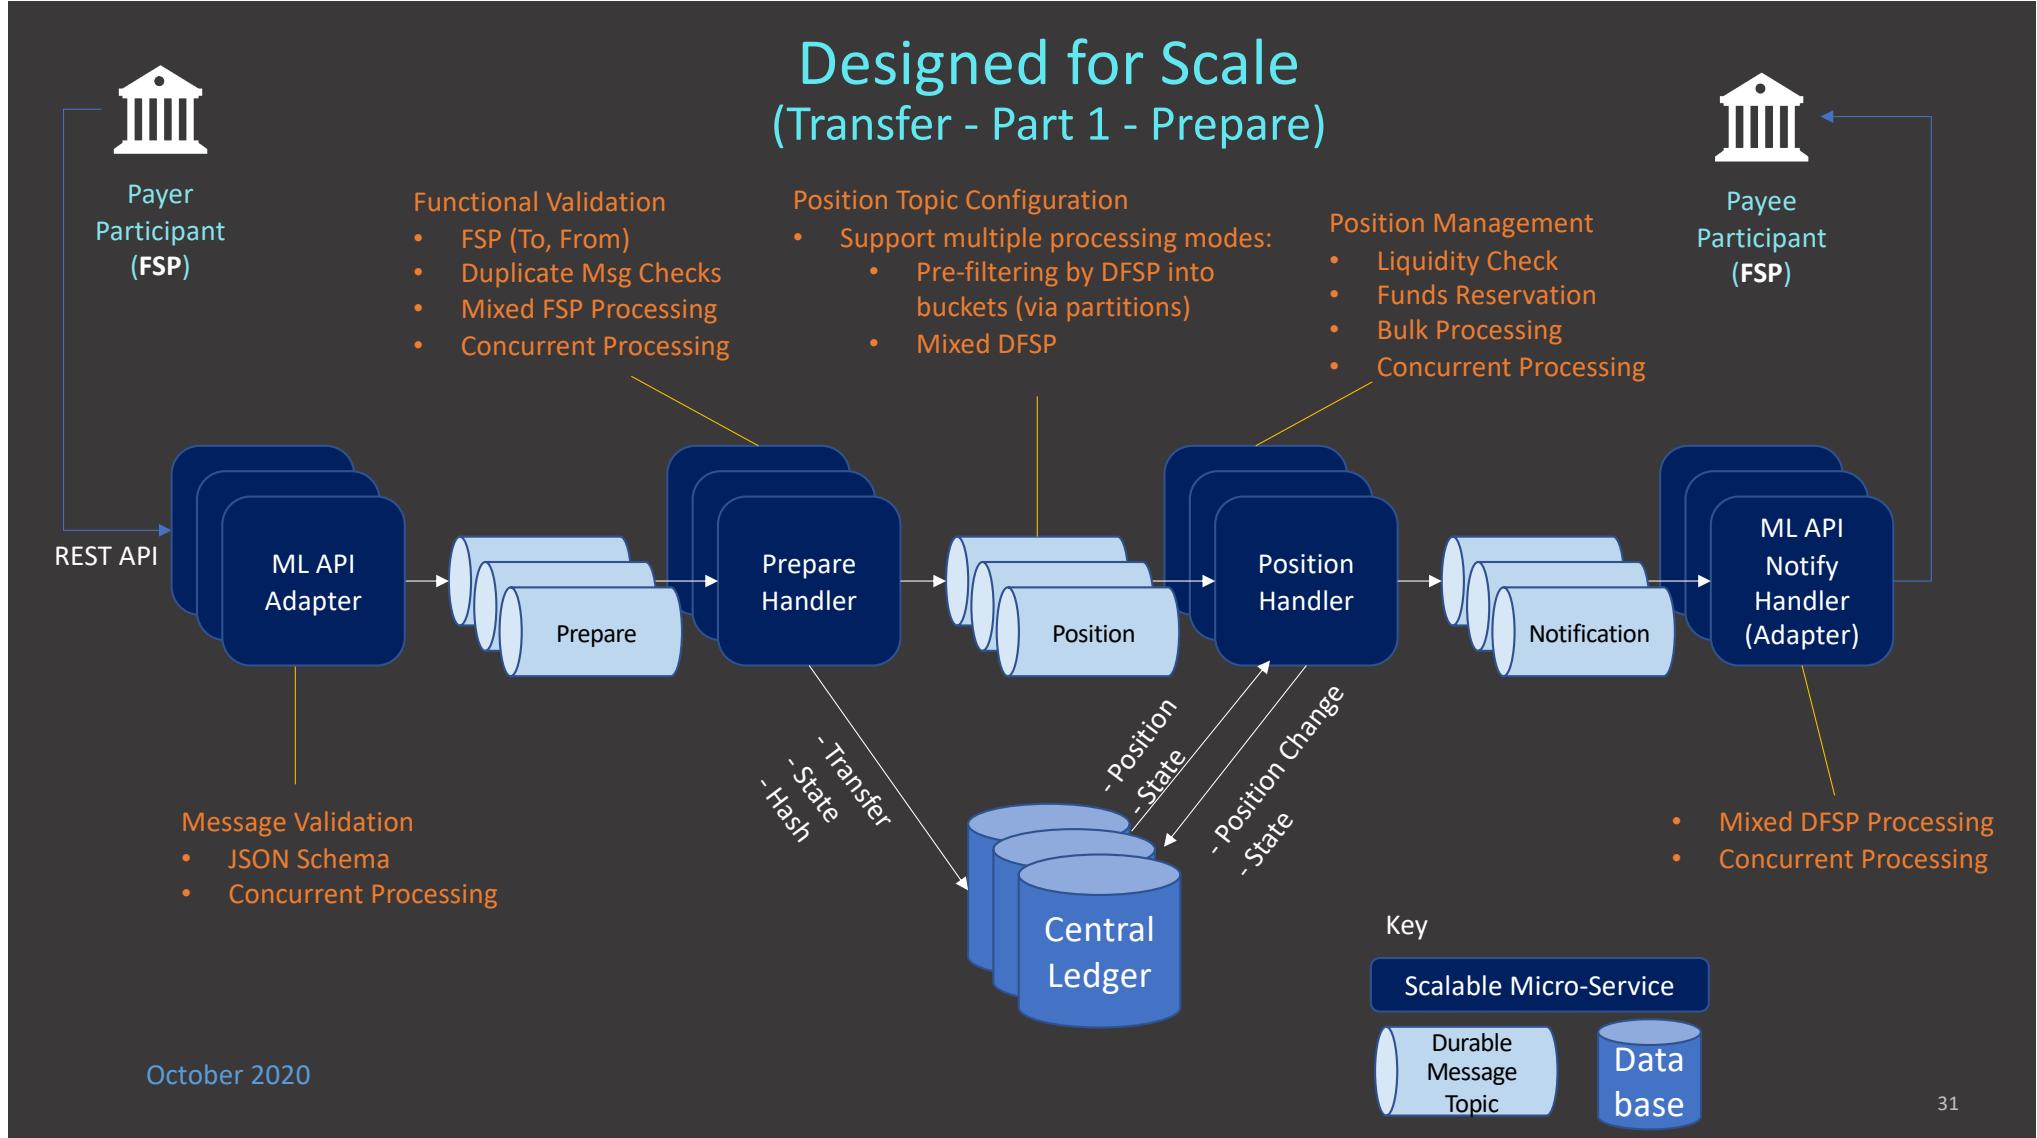


October 2020

* Color-shading of core & supporting components match dependent persistent stores.

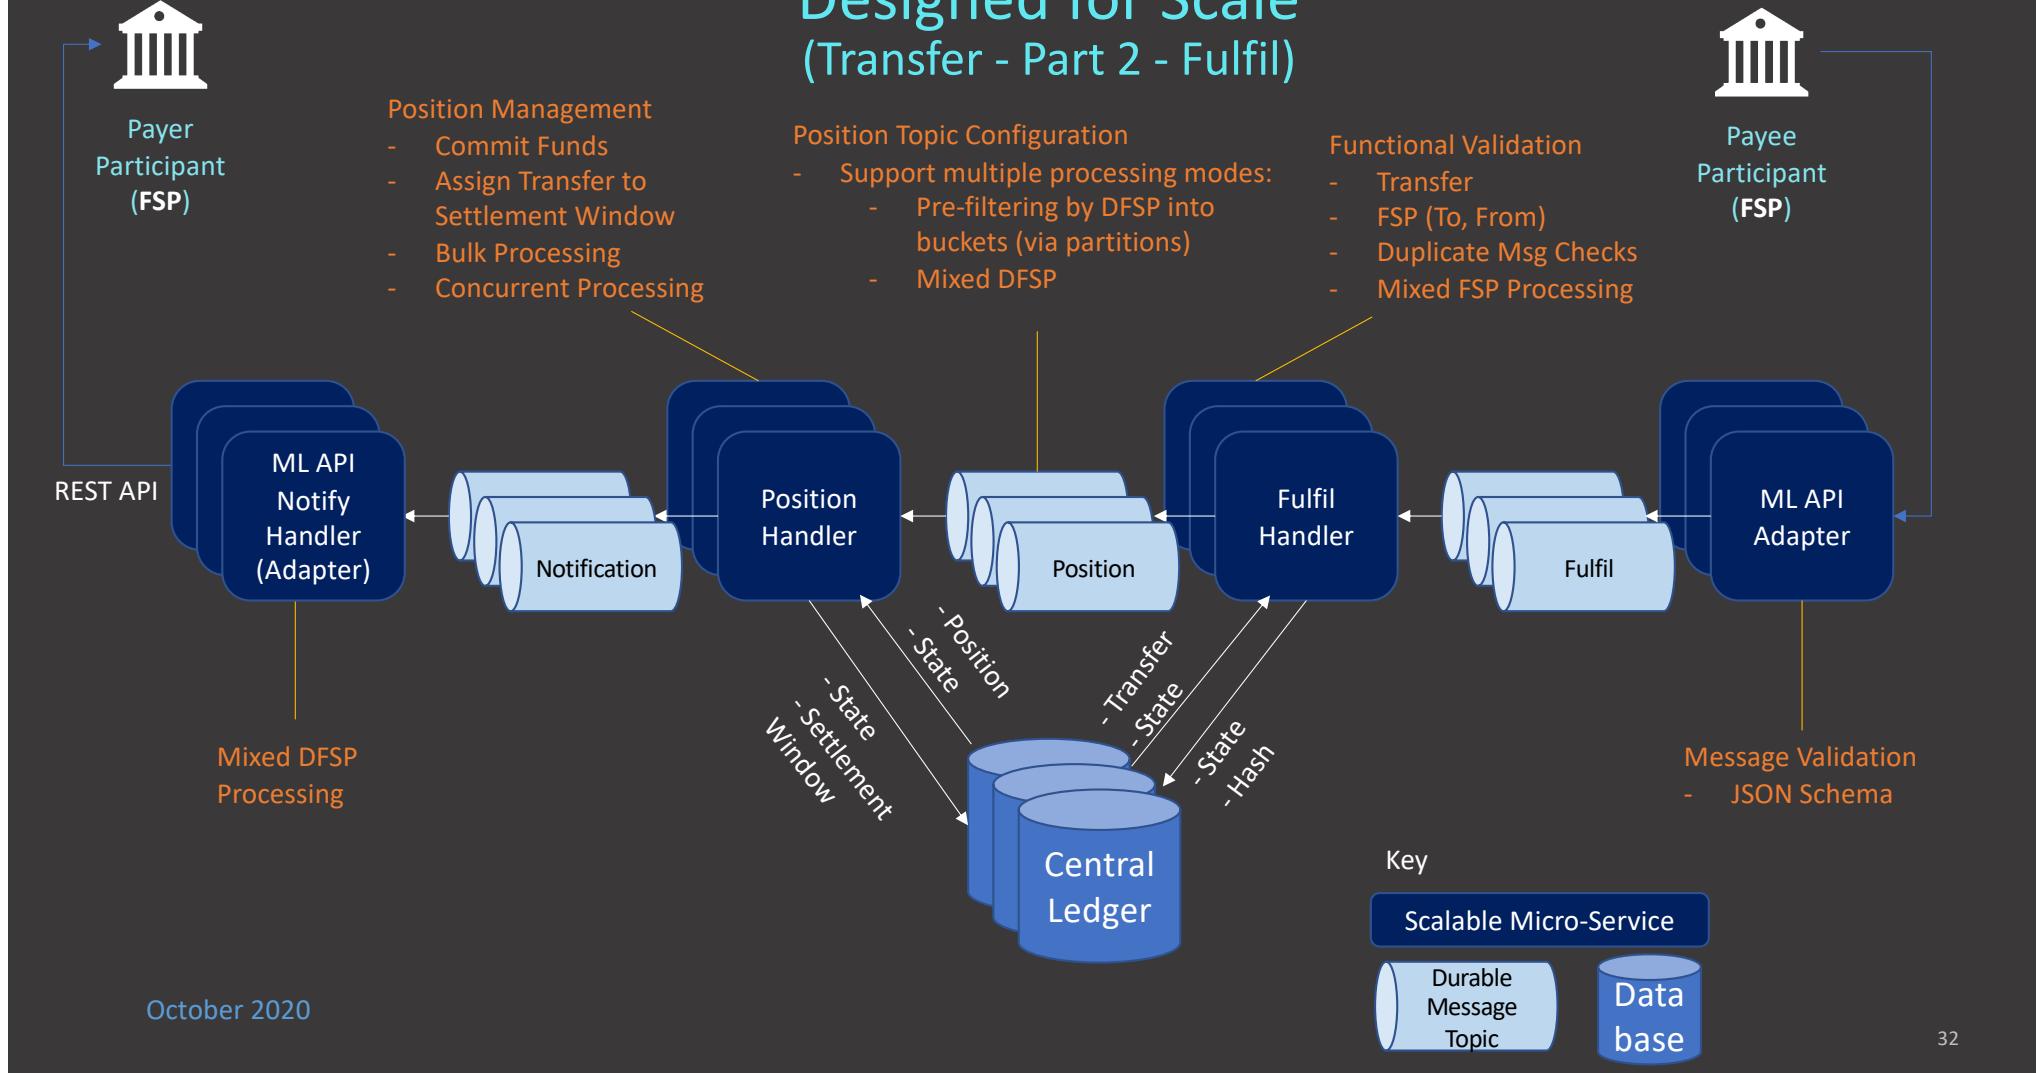
Designed for Scale

(Transfer - Part 1 - Prepare)



October 2020

Designed for Scale (Transfer - Part 2 - Fulfil)

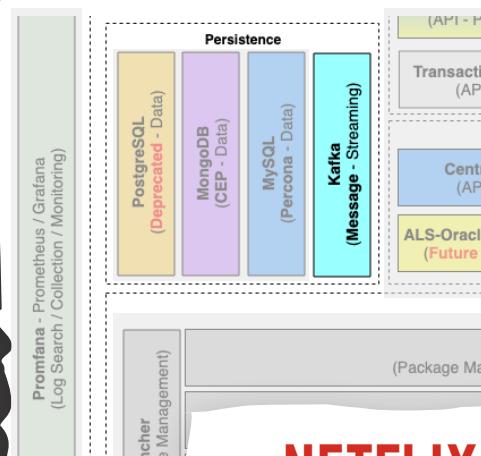


Messaging Platform – Kafka



What is Kafka?

Apache Kafka is a distributed message streaming platform.



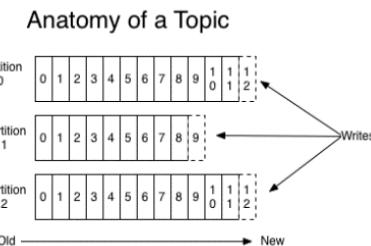
Why Kafka?

Publish and subscribe to streams of records, similar to a message queue or enterprise messaging system.

Store streams of records in a fault-tolerant durable way with history being saved for a desired period.

Ref: <https://kafka.apache.org/>

Process streams of records as they occur.



- For each topic, the Kafka cluster maintains a partitioned log.
- Each partition contains a sequence of records that is
 - Ordered; and
 - Immutable
- The records in the partitions are each assigned a sequential id number called the *offset* that uniquely identifies each record within the partition.

NETFLIX

PayPal ORACLE®

IBM

Rabobank

Used By

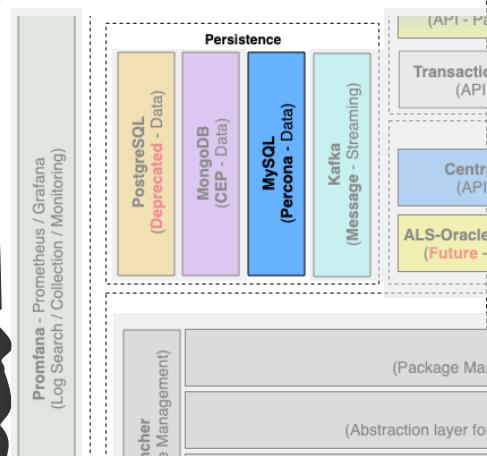
Storage Platform – Percona XtraDB Cluster



PERCONA

What is Percona XtraDB?

An open source, cost-effective, and robust MySQL clustering solution for businesses.



Why Percona XtraDB?

Ref: <https://www.percona.com/software/mysql-database/percona-xtradb-cluster>



Cost-effective HA and scalability for MySQL with both Open Source and Enterprise support options



Increased read/write scalability



Zero data Loss



Multi-master replication



Works on-premises, cloud, hybrid, WAN, LAN, Kubernetes support (Helm)

Used By

October 2020



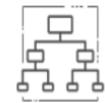
Container Management & Orchestration – Helm

What is Helm?

Open Source Package Manager for Kubernetes through the use of Charts.

Charts help you define, install and upgrade releases for Kubernetes deployment via templates and configuration.

Why Helm?



Manage Complexity

Charts describe even the most complex apps; provide repeatable application installation, and serve as a single point of authority.



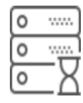
Easy Updates

Take the pain out of updates with in-place upgrades and custom hooks.



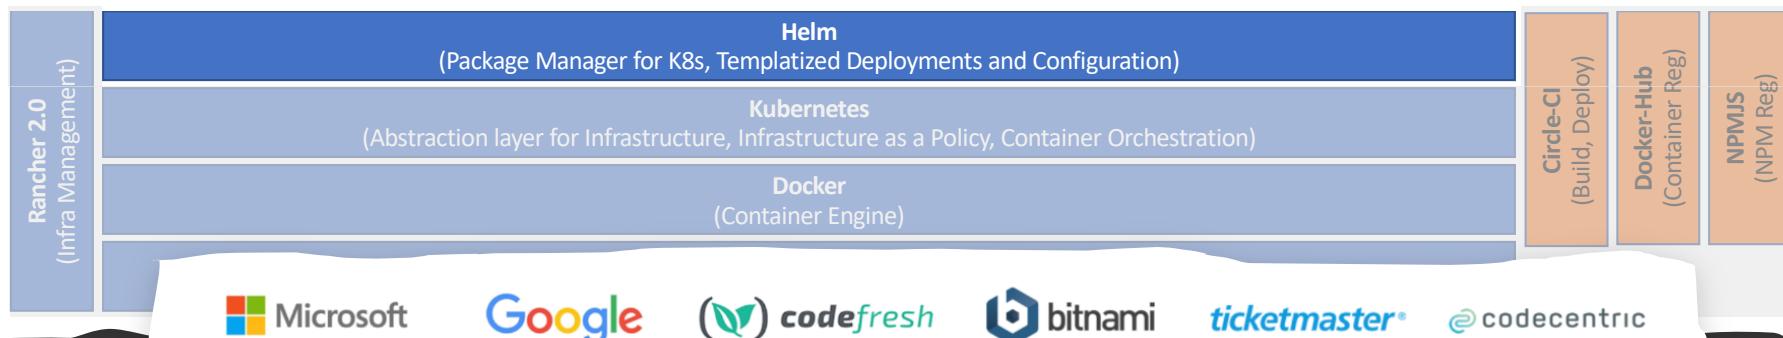
Simple Sharing

Charts are easy to version, share, and host on public or private servers.



Rollbacks

Use `helm rollback` to roll back to an older version of a release with ease.



Container Management & Orchestration – Kubernetes



What is Kubernetes?

Open-source system for automating deployment, scaling, and management of containerized applications.

Why Kubernetes?



Deploy your applications quickly and predictably

- Infrastructure as a Policy
- Abstraction of Infrastructure (Cloud, On-Prem)



Scale your applications on the fly

- Policy rule based scaling
- Limit hardware & resources by scaling horizontally up/down



Roll out new features seamlessly

- Rolling updates



Discoverability

- Dynamic service resolution via DNS



Durability

- Self-healing
- Auto-[placement, restart, replication, scaling] based on Policies
- Load Balancing



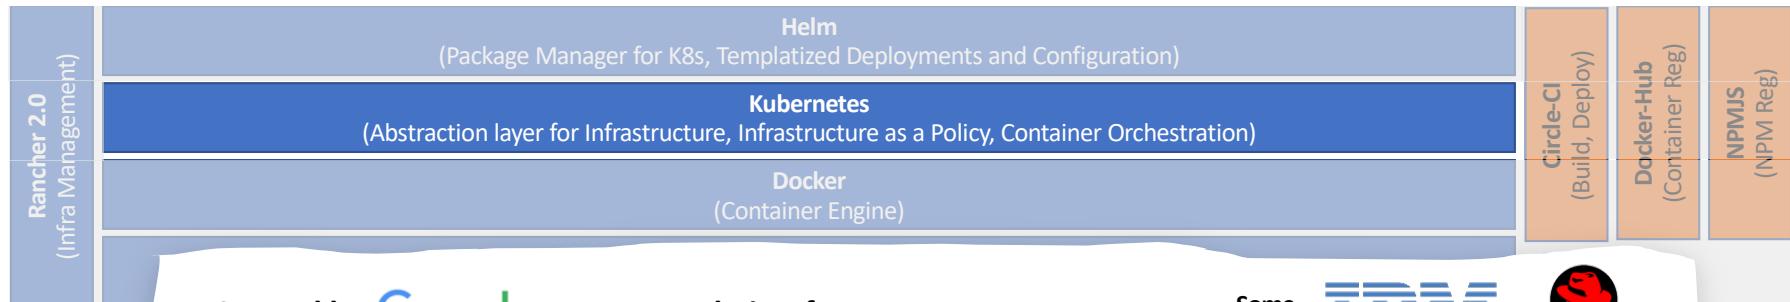
Security

- Isolation through Containers, Network and Namespaces



Operations

- App config & secrets stored in distributed key-value store (etcd)
- Monitoring of containers



Created by **Google** to support their Infrastructure.

Some contributors:





Deployment Architecture – Rancher Overview

Ref: <http://rancher.com>

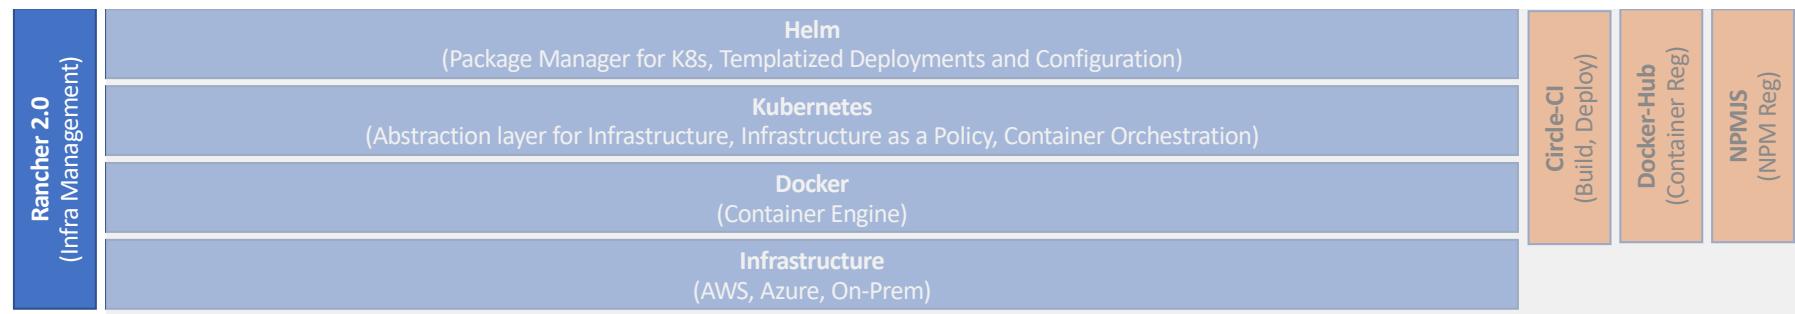
What is Rancher?

Rancher is an enterprise management plane for Kubernetes.

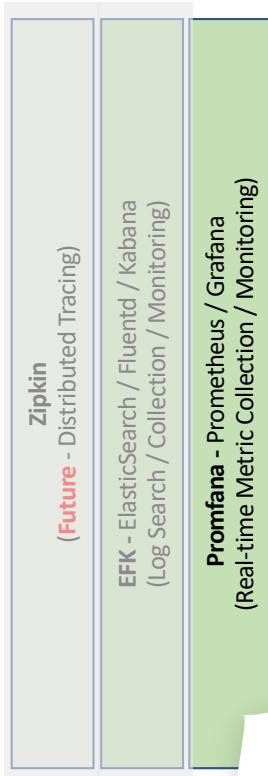
“Every distro. Every cluster. Every cloud.” ~ rancher.com

Why Rancher?

- Kubernetes Management
- Container Management
- Access Management (RBAC)
- Helm Repository Management
- Multi-environment Management (multi k8s clusters, On-prem, Azure, Google, AWS, etc)
- Multi-Provider provisioning (On-prem, Azure, Google, AWS, vSphere)
- Easily scale up/down Kubernetes clusters



Deployment Architecture – Ops Monitoring Overview



What is Metric Instrumentation?

Real-time operational visibility for:

- Performance
- Health
- Alerts

Why Promfana?

- **Metric Instrumentation for Mojaloop**
- **Low overhead on nodejs (histograms + pull metric end-point)**
- **Real-time metric visualization for Performance and Health monitoring of the Mojaloop Stack**

What is Promfana?



Leading open-source instrumentation solution for monitoring



The open platform for beautiful analytics and monitoring

Dimensional data

Prometheus implements a highly dimensional data model. Time series are identified by a metric name and a set of key-value pairs.

Powerful queries

PromQL allows slicing and dicing of collected time series data in order to generate ad-hoc graphs, tables, and alerts.

Great visualization

Prometheus has multiple modes for visualizing data: a built-in expression browser, Grafana integration, and a console template language.

Efficient storage

Prometheus stores time series in memory and on local disk in an efficient custom format. Scaling is achieved by functional sharding and federation.

Simple operation

Each server is independent for reliability, relying only on local storage. Written in Go, all binaries are statically linked and easy to deploy.

Precise alerting

Alerts are defined based on Prometheus's flexible PromQL and maintain dimensional information. An alertmanager handles notifications and silencing.

Many client libraries

Client libraries allow easy instrumentation of services. Over ten languages are supported already and custom libraries are easy to implement.

Many integrations

Existing exporters allow bridging of third-party data into Prometheus. Examples: system statistics, as well as Docker, HAProxy, StatsD, and JMX metrics.



CoreOS



docker



DigitalOcean



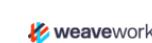
SOUNDCLLOUD



PERCONA

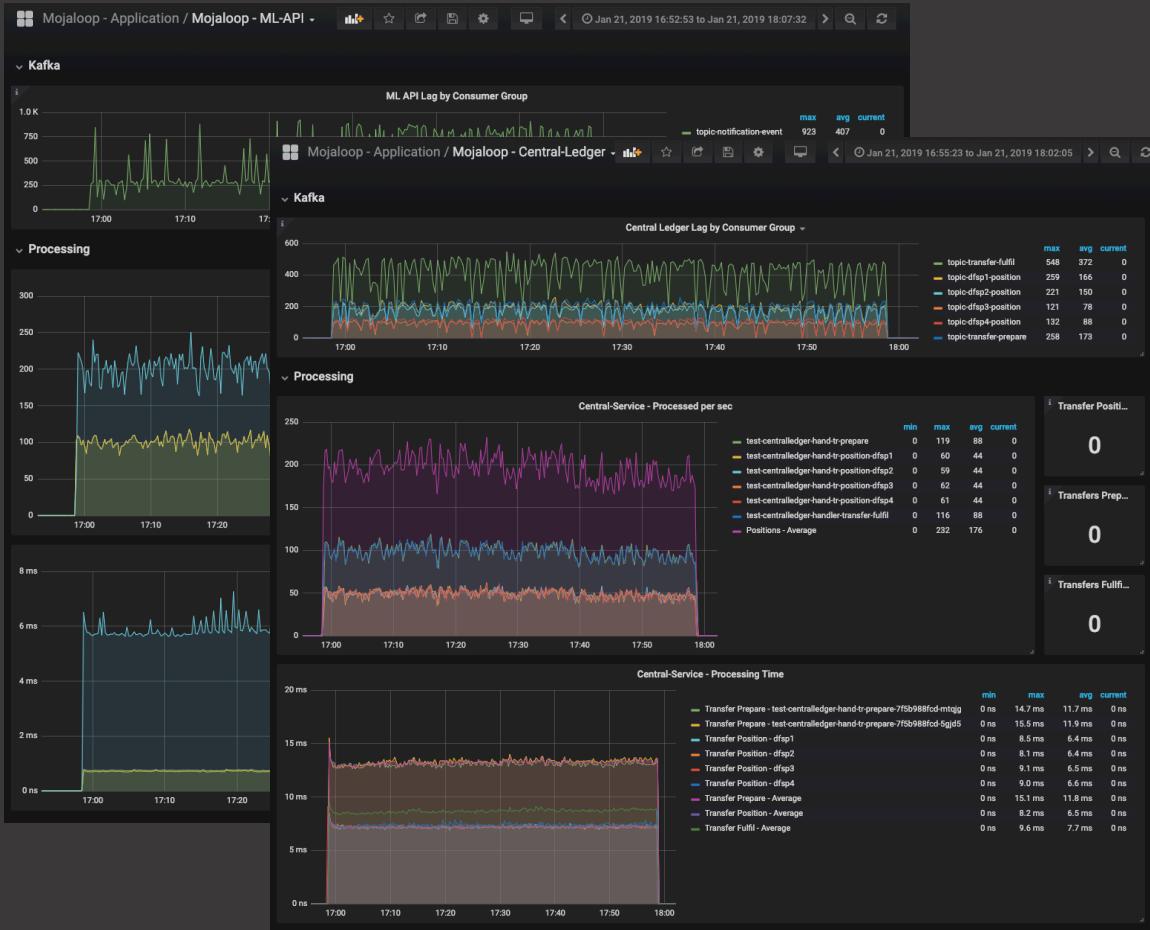


ERICSSON



weaveworks

Deployment Architecture – Ops Monitoring Dashboards



October 2020

Documentation

<http://mojaloop.io/helm/monitoring/>

Mojaloop Application

- ML-API-Adapter -- nodejs + application
- Central-Ledger -- nodejs + application
- Simulators -- nodejs + application

Data Store

- MySQL Overview
- PXC Galera Overview
- PXC Galera Graphs

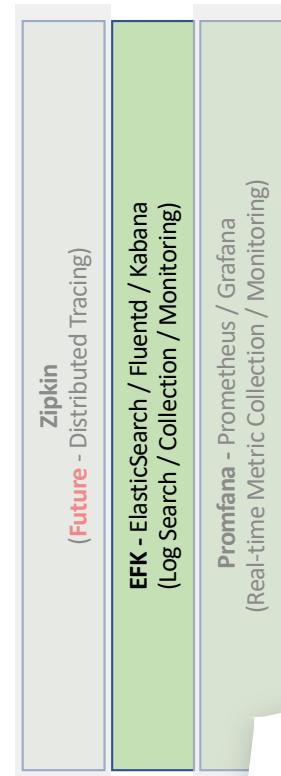
Messaging

- Kafka Cluster Overview
- Kafka Topic Overview

Kubernetes

- Clusters
- Deployments

Deployment Architecture – Ops Logging Overview



What is EFK (aka ELK)?



elasticsearch

Elasticsearch is a distributed, RESTful search and analytics engine.



fluentd

Open source data collector for unified logging layer, with ingestions into Elasticsearch.



kibana

Kibana lets you visualize your Elasticsearch data and navigate the Elastic Stack.

Why EFK?

- Central location and storage of all Mojaloop log files
- Management of log data (persistence, long-term storage, etc)
- Management of alert/events based on log data
- Log files are indexed and searchable
- Assist with tracing & trouble shooting Mojaloop's distributed micro-service logs

Why E-F-K and not E-L-K?

- Fluentd is used instead of Logstash due to its support & seamless integration for Kubernetes (k8s).
- K8s Pods/Containers are easily collected by Fluentd and ingested into Elasticsearch using the underlying K8s logging architecture.





Deployment Architecture – CircleCI Overview

What is CircleCI? Cloud based Continuous Integration & Deployment Platform

Ref: <http://circleci.com>

VCS Integration

CircleCI integrates with GitHub, GitHub Enterprise, and Bitbucket. Every time you commit code, CircleCI creates a build.

Automated Testing

CircleCI automatically tests your build in a clean container or virtual machine.

Automated Deployment

Passing builds are deployed to various environments so your product goes to market faster.

Notifications

Your team is notified if a build fails so issues can be fixed quickly.

Why CircleCI?



Workflows for Job Orchestration

Orchestrate customizable job execution (such as build, test, deploy), giving complete control over your development process.



Language-Agnostic Support

Supports any language that builds on Linux or macOS, including C++, Javascript, .NET, PHP, Python, and Ruby.



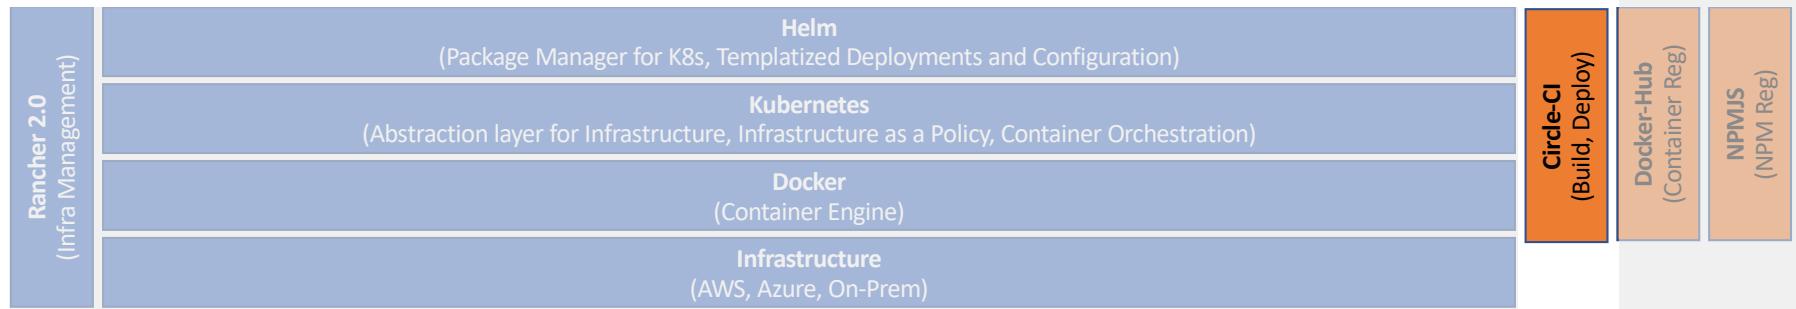
First-Class Docker Support

Run any image from Docker's public/private registry or other common registries. Build Docker images, access Docker layer caching, Compose.



Powerful Caching

Speed up builds with expanded caching options, including images, source code, dependencies, and custom caches. Full control over cache save and restore points for optimal performance.



* Forrester names CircleCI a leader (<https://www2.circleci.com/circleci-forrester-wave-leader-2017.html>)

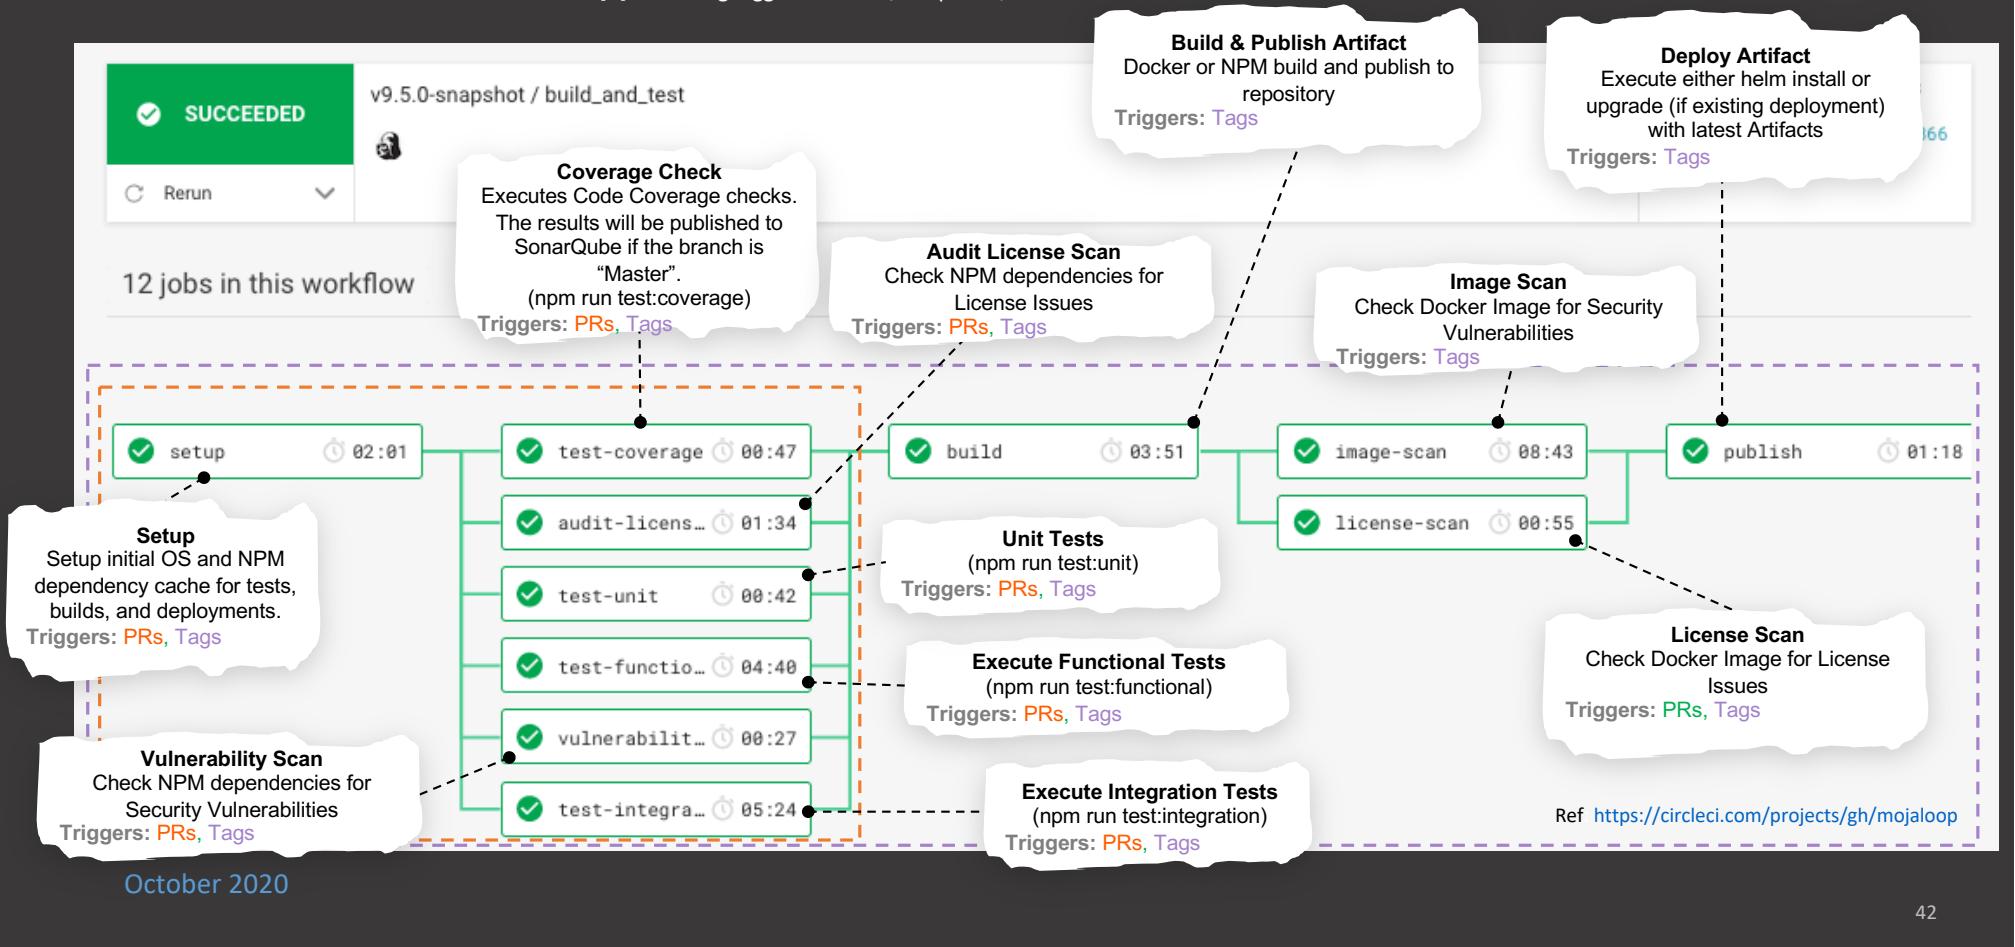
October 2020



Deployment Architecture – CI/CD Pipeline

Triggers:

- [●] Pull-Requests (inc. from external Forks)
- [●] Publishing tagged Releases / Snapshots, etc





Documentation – GitBooks Overview

What is Gitbooks?

An open-source open documentation framework where teams can document everything from products, to APIs and internal knowledge-bases based on open-standards with community driven plugins.

Why Gitbooks?



Markdown

Lightweight markup language with plain text formatting syntax supporting standard HTML, and CSS.



Imbed Generated Content

Embed generated sequence diagrams, openapi/swagger docs, etc.



Search

Find what you are looking for.



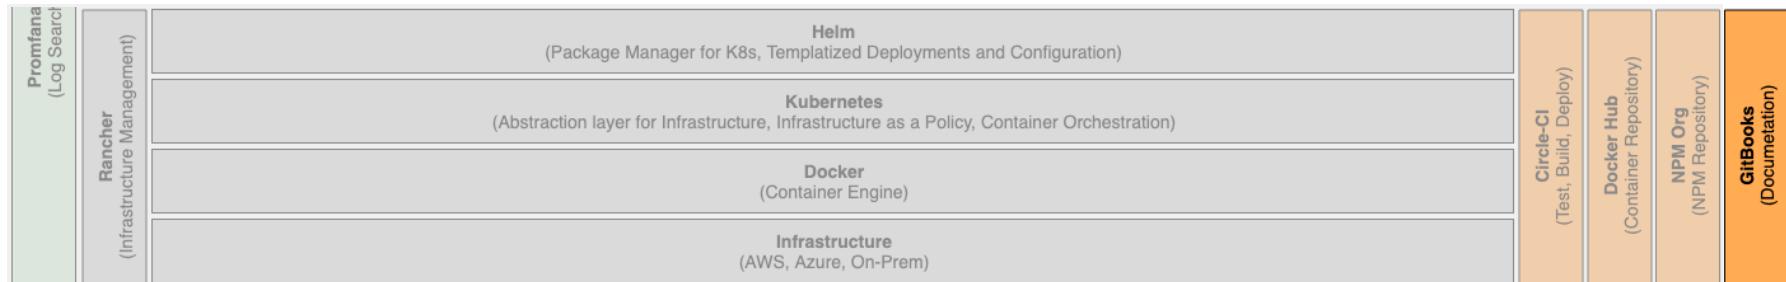
Plugins

Community plugins for generating content (e.g. plantuml, openapi/swagger docs), providing integration to Github, Slack, etc and themes (e.g. ToCs, Navigation, etc)



Cli

Gitbook-cli to build static-content, with support for local testing. Also supports auto-build sense when changes are made locally when testing.



Mojaloop

Technical Overview

1. What is Mojaloop?
2. High-level Overview, Static Demo & Architecture
3. Component Architecture
4. Global foot-print & Roll-out

Mojaloop's Global Foot-print





Mojaloop Questions?

October 2020





Mojaloop

Appendix

October 2020



Sequence Diagrams: Overview

1. Overview of Sequence Diagrams
2. High level Design for :
 - a. Prepares, Positions, Fulfils, Rejections, Timeouts
 - b. Operations: Settlements, Positions, Limits, callback URLs, etc.
3. List of Sequence Diagrams
4. Location: <https://mojaloop.io/documentation/mojaloop-technical-overview>

API First Approach

Key

- [●] Implemented using API First Approach
- [●] Not implemented – To be re-factored in future

Background

Legacy Mojaloop API components were implemented using a Code First approach (even if a contract existed).

What is API First Approach?

An Interface Contract is defined up-front, which is then used to generate Base-code which includes validations, routes, documentation, etc.

Benefits?

- Faster implementation time as Base-code is auto-magically created for developers, only requiring logic to be implemented.
- QA Improvement
 - Validations are done auto-magically adhering to contract
 - Routes adhere to contract
 - Documentation adheres to contract

Mojaloop & Operational APIs

- [●] Mojaloop-API-Adapter
- [●] Central-Ledger
- [●] Central-Settlements

Example:

<https://github.com/mojaloop/central-settlement>



Tools used

Hapi-openapi

- API schema validation.
- Routes based on the OpenAPI document.
- API documentation route.
- Input validation.
- Ref: <https://github.com/krakenjs/hapi-openapi>

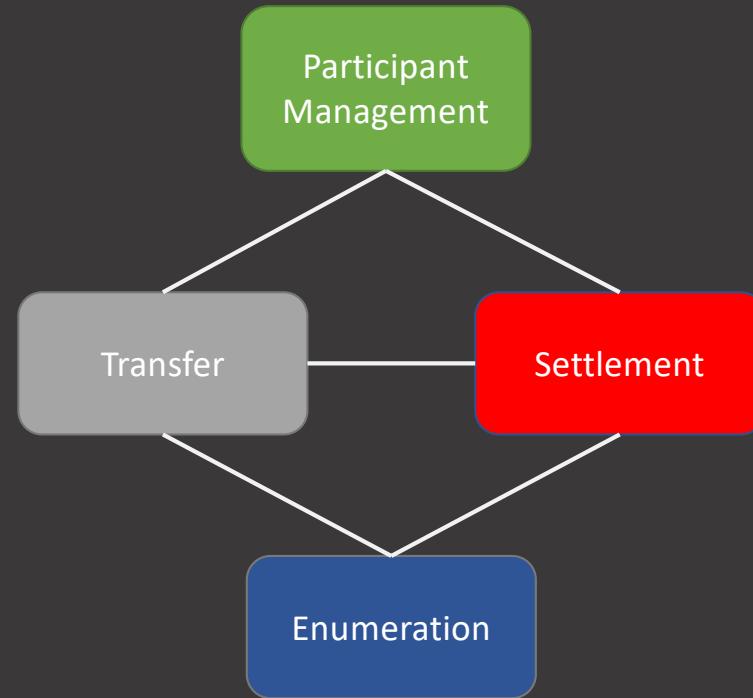
Database Design

Database Design Objectives

1. Align the data model to support to the Mojaloop specification
2. Design for efficiency and reduce locking
3. Data integrity maintained through a idempotency pattern

How was this achieved?

1. Redesign scheme to reflect the Mojaloop API and support the state transitions.
2. Inserts are used instead of updates to ensure reduced locking, and increased speed of data “updates”
3. Inserts are used to ensure an idempotent pattern for data “updates” is maintained. The result being that all “updates” have a history.
4. Support for batch processing of DFSP position, maintain repeatable and deterministic outcomes of rule processing and “running balance” for positions.



Optimization – Position Handler Processing Modes



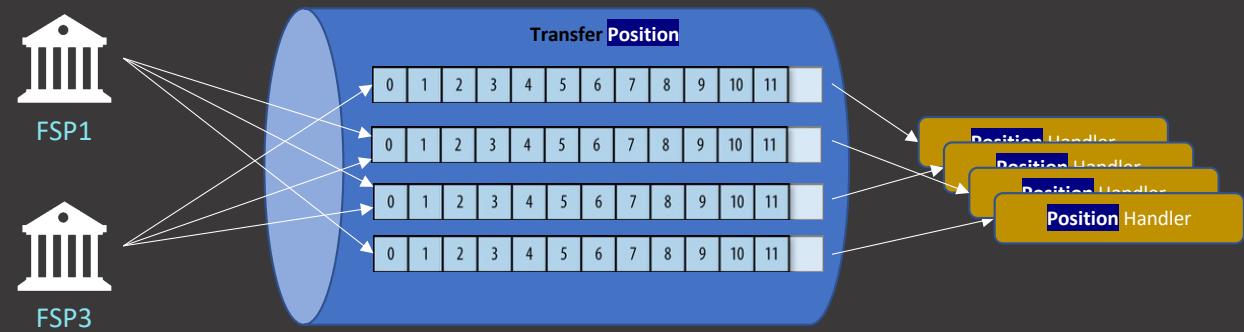
Pre-filtering of DFSPs into Buckets

- Partition strategy is keyed for Positions.
- Message keys are used to route each message to a specific Partition (Bucket).
- Dedicated handlers for DFSP specific workloads (Future: Easier In-memory processing)
- *Possible issue/limitation*:
 - Limited scalability by number of FSPs
 - Possible idle handlers (cost implication)
 - Requires process for increasing partitions on Kafka (possible downtime, migration strategy, or temporal performance hit).



Mixed DFSP Processing

- Partition strategy can be configured as Keyed (same as PI4) or Random
- Improved scalability (no limit)
- All Handlers are utilized (efficient hardware usage)
- Future possibility of consolidating handlers (prepare, position & fulfil)
- *Possible issue/limitation*:
 - Increased contention & locking on the database persistence layer.



Definition of terms: 7. participant position

1. A *participant position* is a net position of all activity
2. A position applies to a switch account for a scheme participant in a specific currency.
3. A position represents:
 1. The activity in an account which is a consequence of transfers flowing through the system
 2. Adjustments to the account as directed by the scheme
4. The position is calculated as:
 1. The sum of reservations made for prepared but unfulfilled transfers where the participant is the payer.
 2. The sum of fulfilled transfers where the participant is the payee.
 3. Any changes to the account position directed by the scheme.

Definition of terms: 7. Net Debit Cap

1. In principle, a participant is not allowed to trade in the system if their [position](#) exceeds their [Net Debit Cap Limit](#).
 1. Which means in this context: *if the sum of credits to and debits from their switch account is less than the net of their transfer activities in the account.*
2. The difference between the sum of credits and debits to a participant's switch account and the net of a participant's transfer activities in the account is their *liquidity cover*.
 - ❖ In this case, “transfer activities” includes transfers where the participant is the payer and the transfer has been requested but not yet completed.
3. So the principle is: a participant must always be able to cover their liabilities. We call this a *liquidity cover requirement*
4. Seems clear enough...